

PRISM: Streamlined Packet Processing for Containers with Flow Prioritization

(ICDCS '22)

Manish Munikar¹, Jiabin Lei², Hui Lu², Jia Rao¹

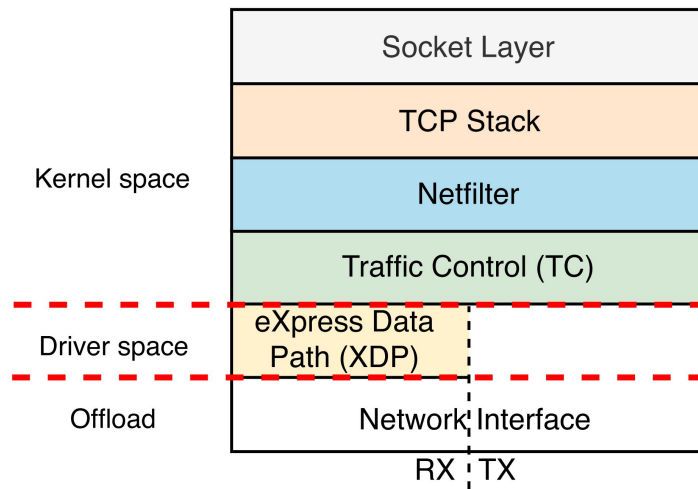
¹ *University of Texas at Arlington*

² *Binghamton University*




Kernel network stack has been very successful ...

- Easy **abstraction** for network I/O
 - Generic sockets interface
- Flexible and **modular**
- Open-source, **stable**, **secure** and **reliable**
- Designed to provide overall efficiency
- Successful for more than **three decades**!



... until recently?

- Network devices are getting faster *faster*
 - 10 Gbps (2002) → 200 Gbps (2017)
- CPU speed is stagnating at 2–3 GHz
 - Scaling *horizontally* instead
- Kernel network stack has been trying to keep up...
 - Packet steering to support multi-core
 - NAPI for dynamic interrupt/poll mode
 - Packet coalescing / batching
 - Checksum offloading
- **New bottleneck** for high-performance network applications



Container networks

- Containers are revolutionizing cloud
 - Lightweight OS-level virtualization
 - Higher consolidation density
 - Faster and easier to manage
- Communicate using overlay network
 - VXLAN encapsulation



Google Cloud



docker



kubernetes

Container networks

- Containers are revolutionizing cloud
 - Lightweight OS-level virtualization
 - Higher consolidation density
 - Faster and easier to manage
- Communicate using overlay network
 - VXLAN encapsulation




Google Cloud



docker




kubernetes



Container networks

- Containers are revolutionizing cloud
 - Lightweight OS-level virtualization
 - Higher consolidation density
 - Faster and easier to manage
- Communicate using overlay network
 - VXLAN encapsulation



Cloud application types

Cloud application types

- Bulk data transfer
 - File transfer, backup, sync
 - Video streaming
 - Many/large packets
 - **Throughput-sensitive**



Cloud application types


- Bulk data transfer
 - File transfer, backup, sync
 - Video streaming
 - Many/large packets
 - **Throughput-sensitive**
- Request-Response
 - Web / database request, RPC
 - Control plane signals
 - Few/small packets
 - **Latency-sensitive**





{ **REST:API** }



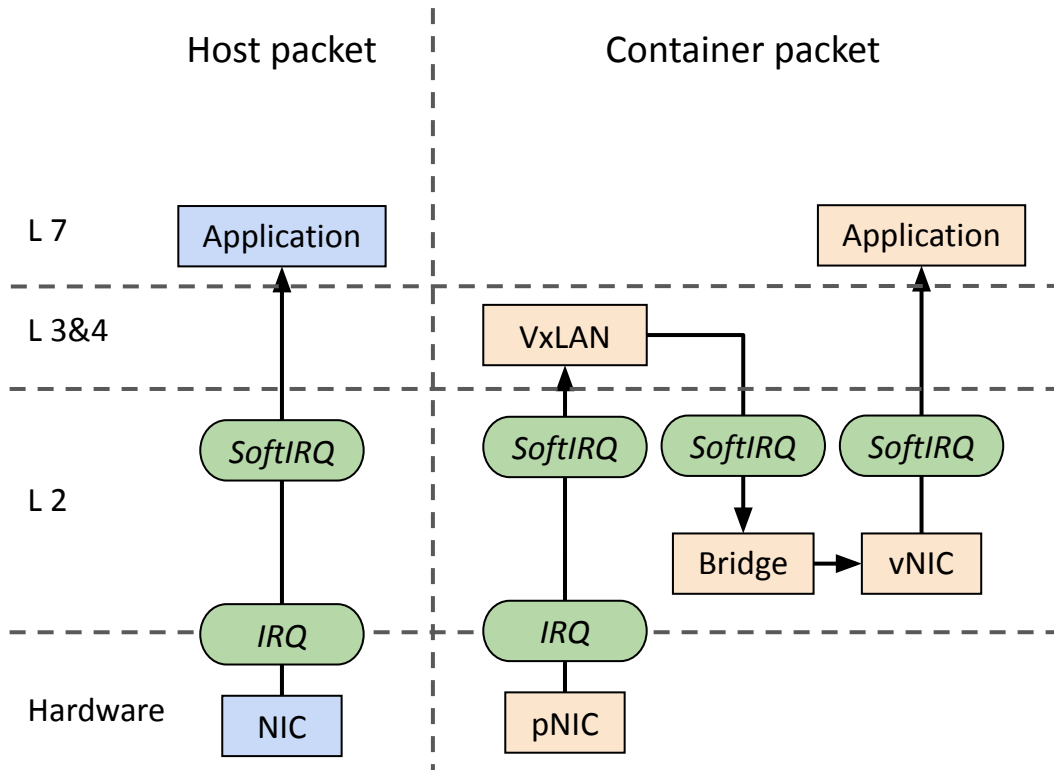

Container packet processing



Container packet processing



Container packet processing



Related works

Related works

- Kernel stack optimizations [RPS, RSS, RFS, FlexSC, ...]
 - System call batching
 - Packet steering / load balancing
 - Interrupt coalescing
 - Zero-copy

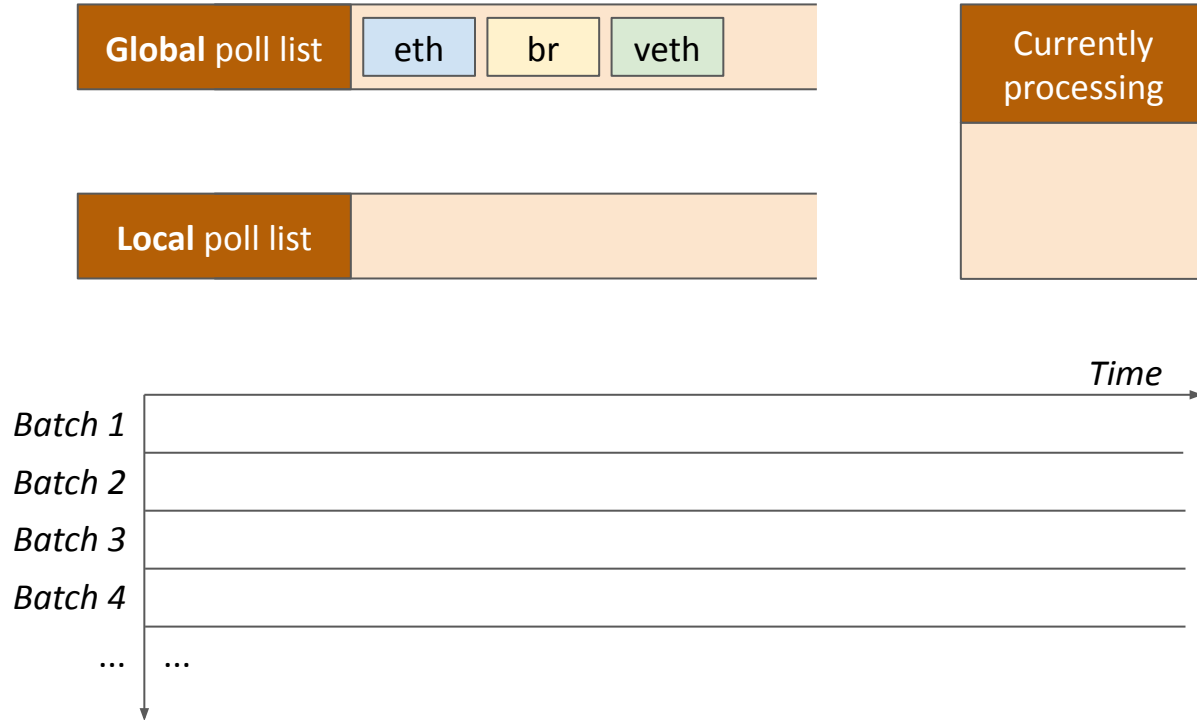
Related works

- Kernel stack optimizations [RPS, RSS, RFS, FlexSC, ...]
 - System call batching
 - Packet steering / load balancing
 - Interrupt coalescing
 - Zero-copy
- Kernel bypass [DPDK, NetMap, FD.io, XDP, Mellanox ASAP², ...]
 - Avoids most kernel overheads
 - Dedicates core for packet polling and uses custom network stack
 - Hardware offload

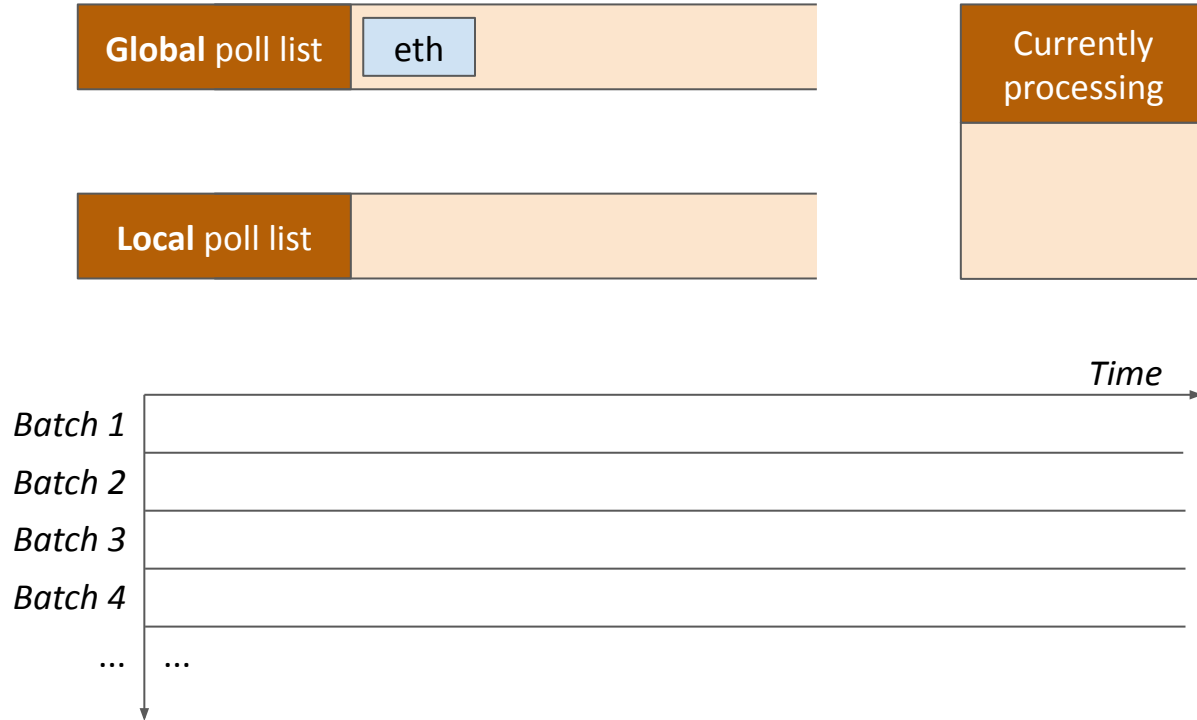
Related works

- Kernel stack optimizations [RPS, RSS, RFS, FlexSC, ...]
 - System call batching
 - Packet steering / load balancing
 - Interrupt coalescing
 - Zero-copy
- Kernel bypass [DPDK, NetMap, FD.io, XDP, Mellanox ASAP², ...]
 - Avoids most kernel overheads
 - Dedicates core for packet polling and uses custom network stack
 - Hardware offload
- Container network optimization [Slim, FALCON, FreeFlow, ...]
 - Connection metadata manipulation
 - More fine-grained (flow-device level) packet steering

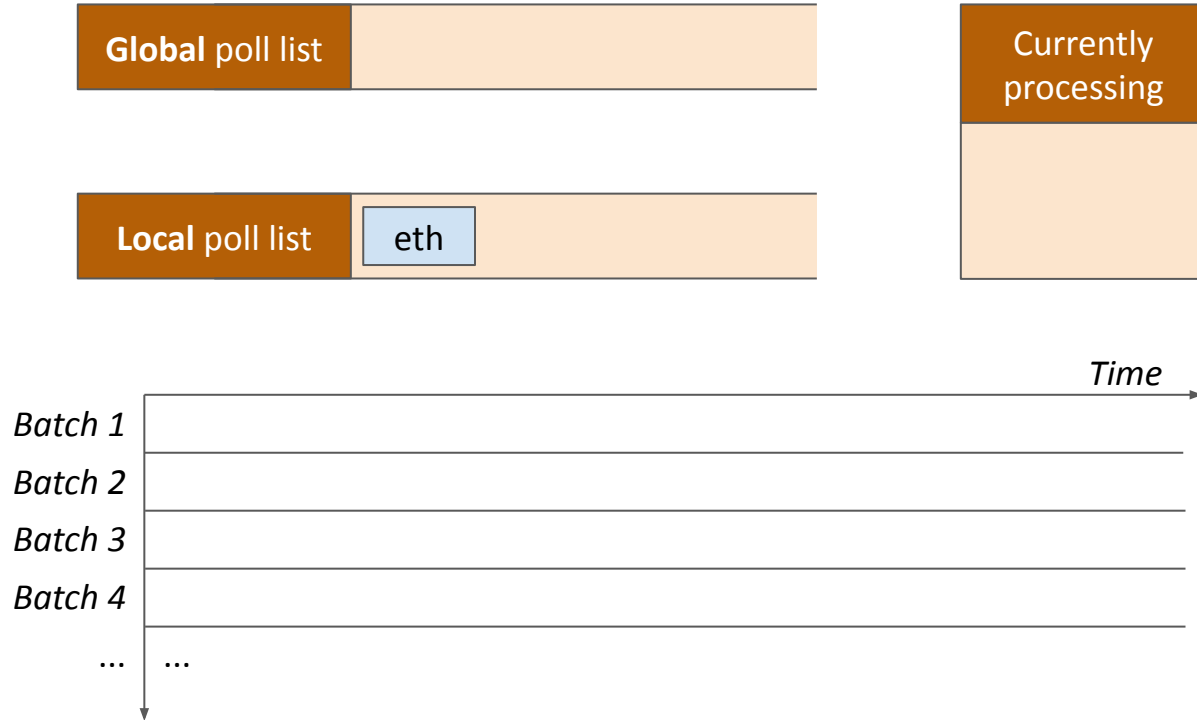
Interleaved NAPI device processing



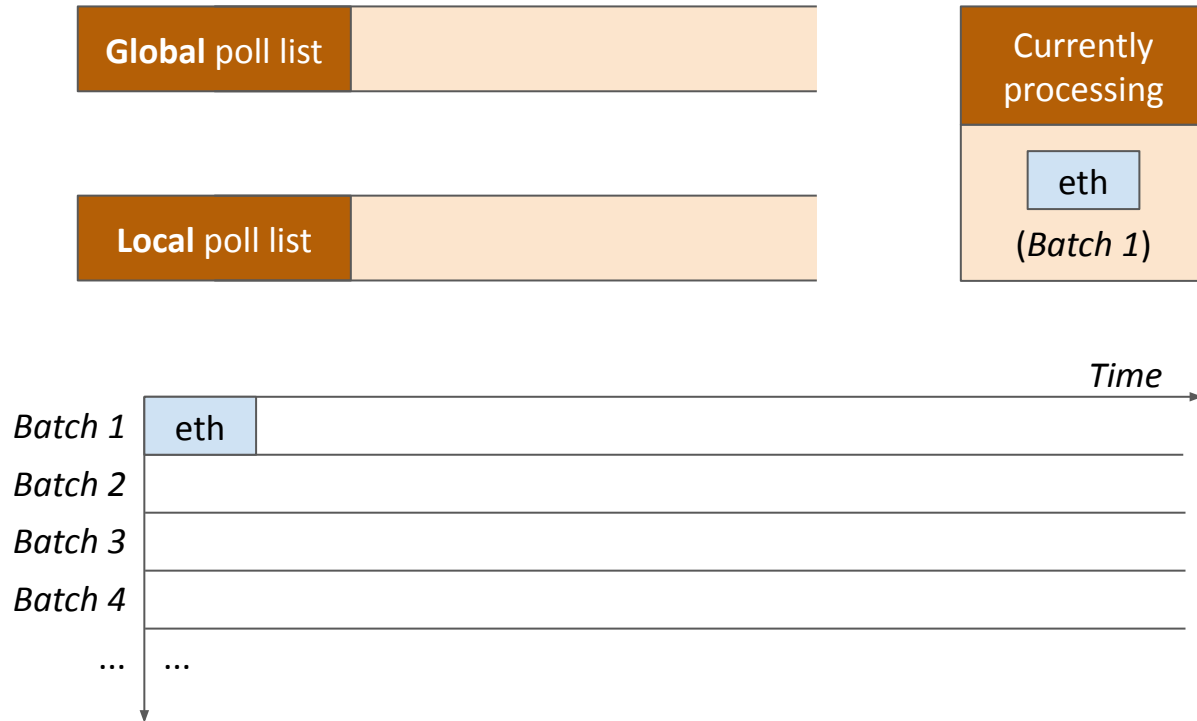
Interleaved NAPI device processing



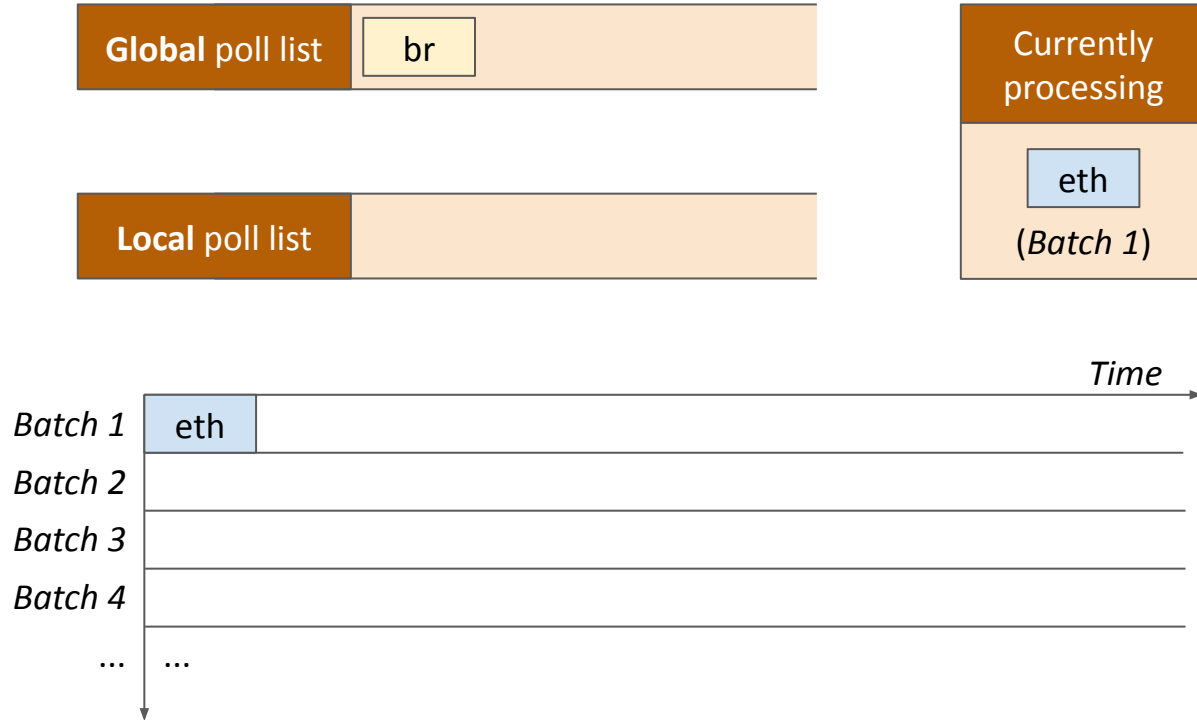
Interleaved NAPI device processing



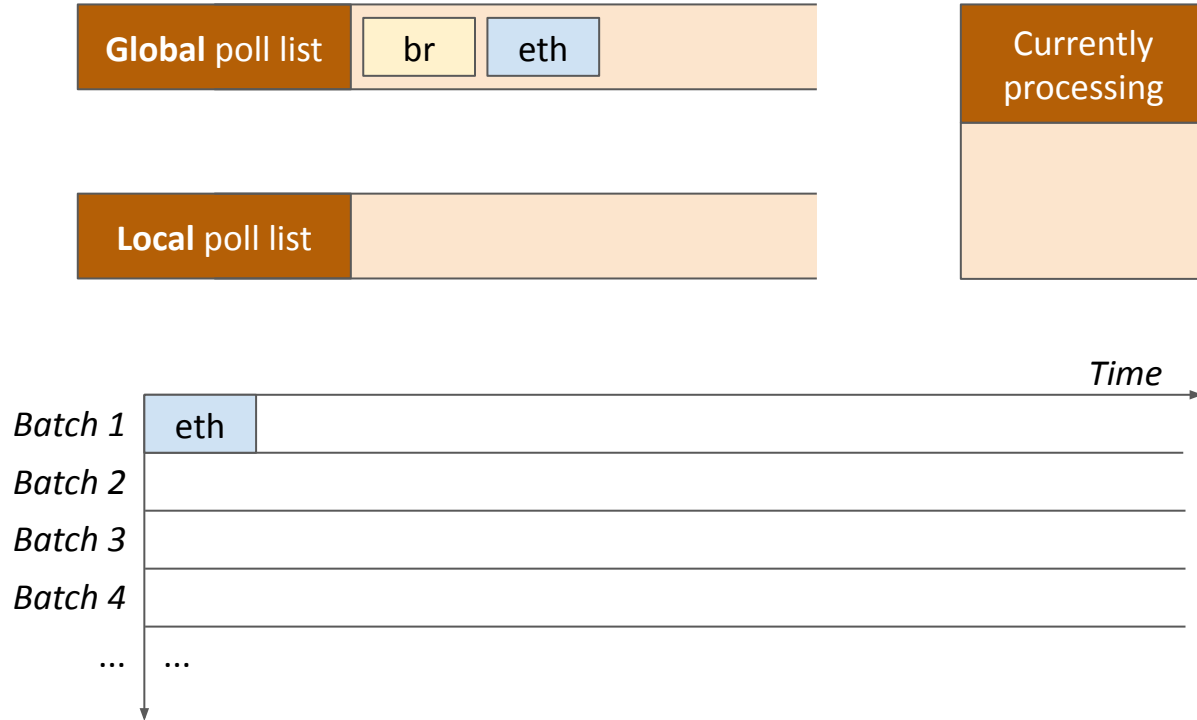
Interleaved NAPI device processing



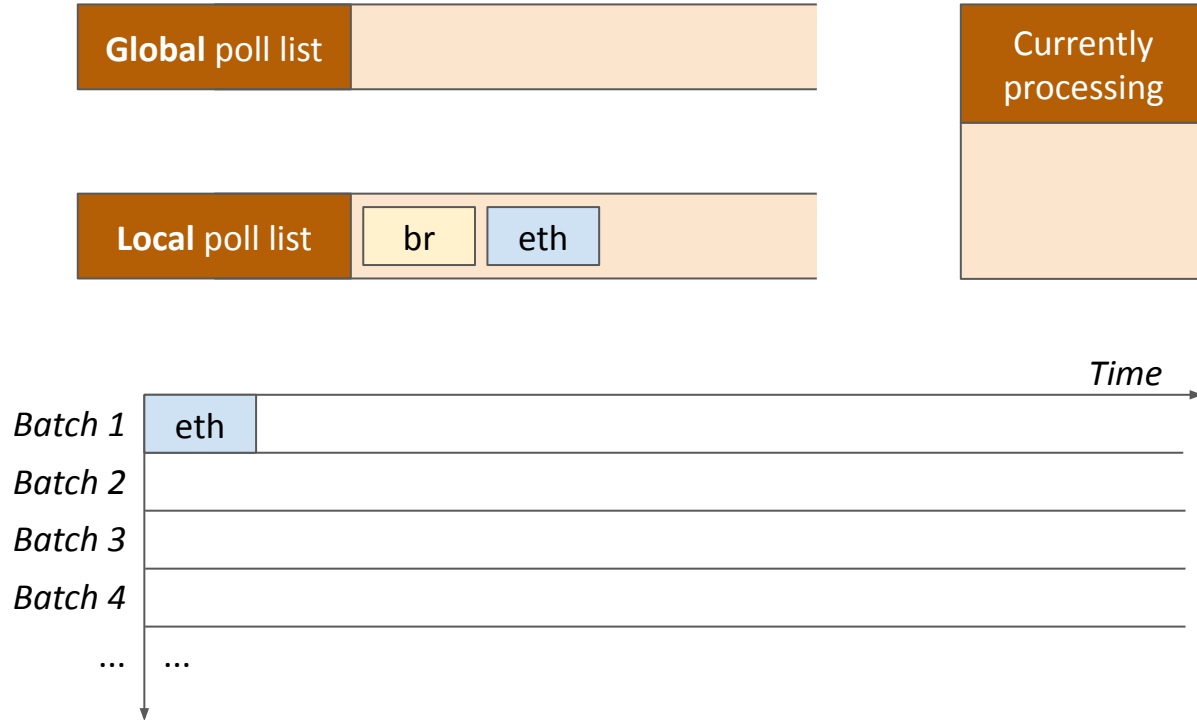
Interleaved NAPI device processing



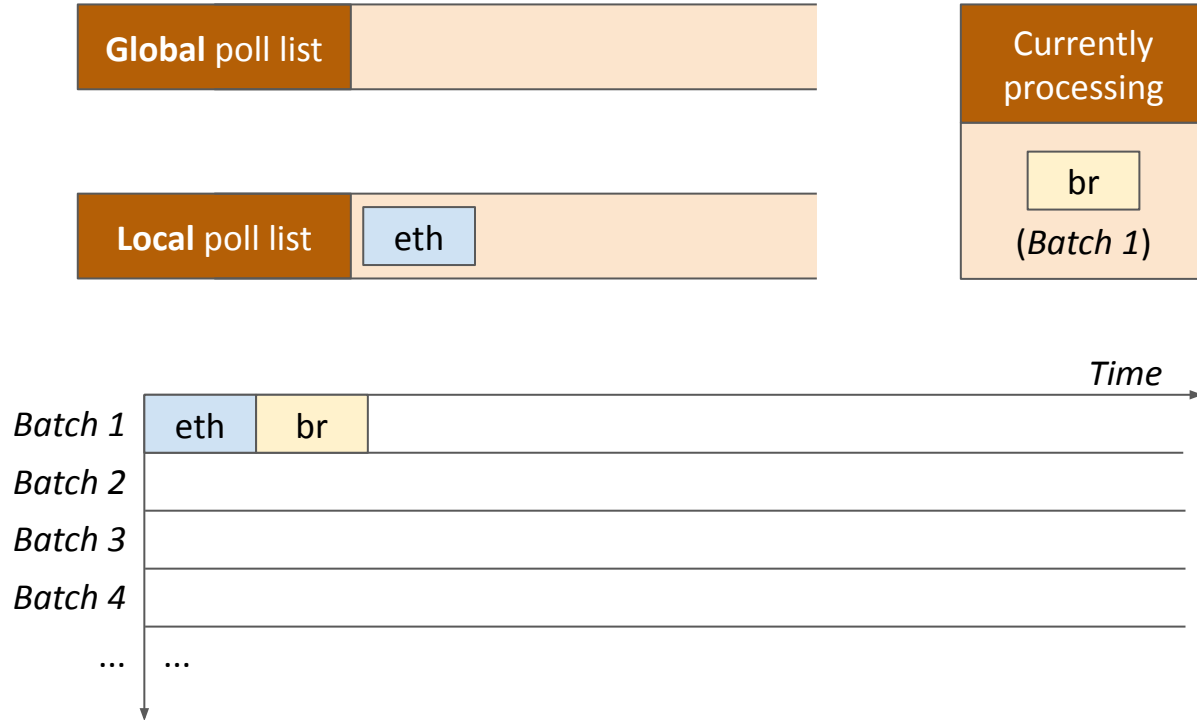
Interleaved NAPI device processing



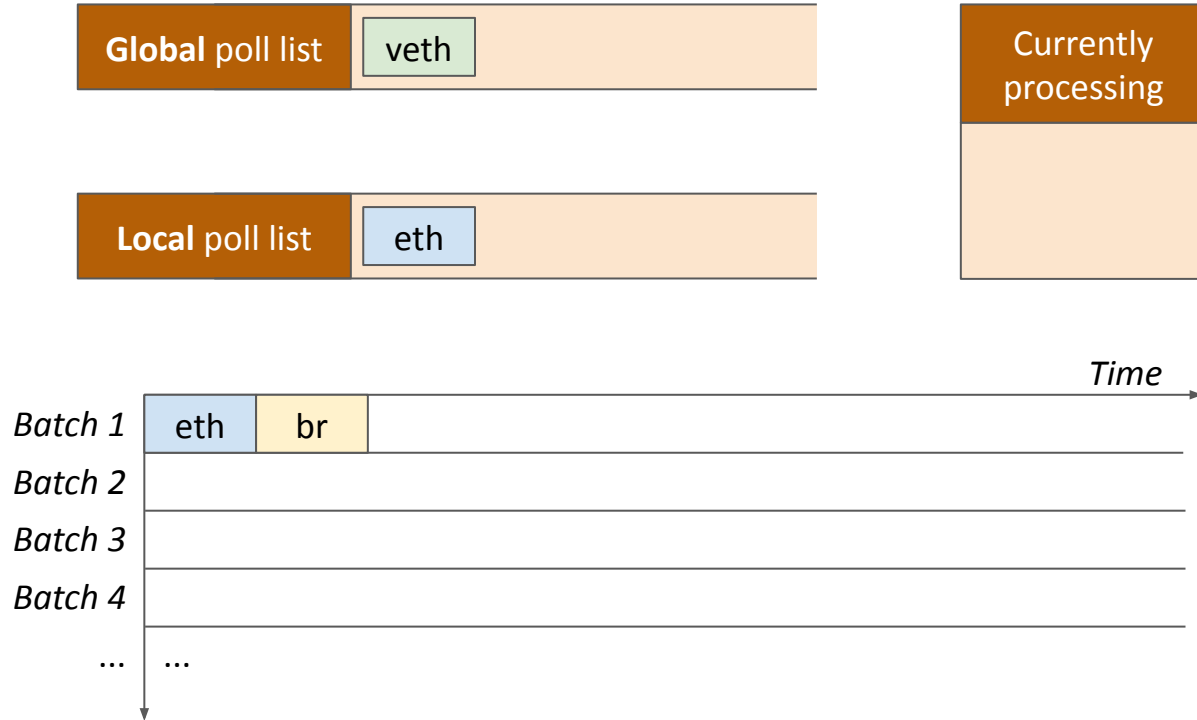
Interleaved NAPI device processing



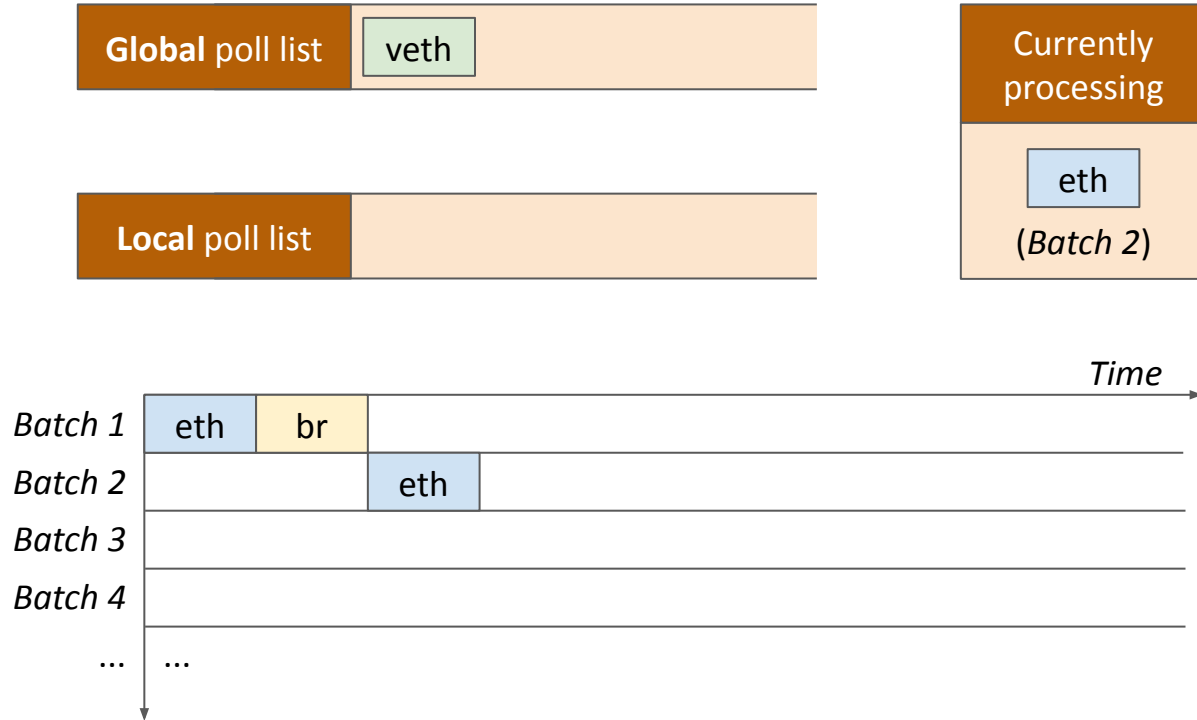
Interleaved NAPI device processing



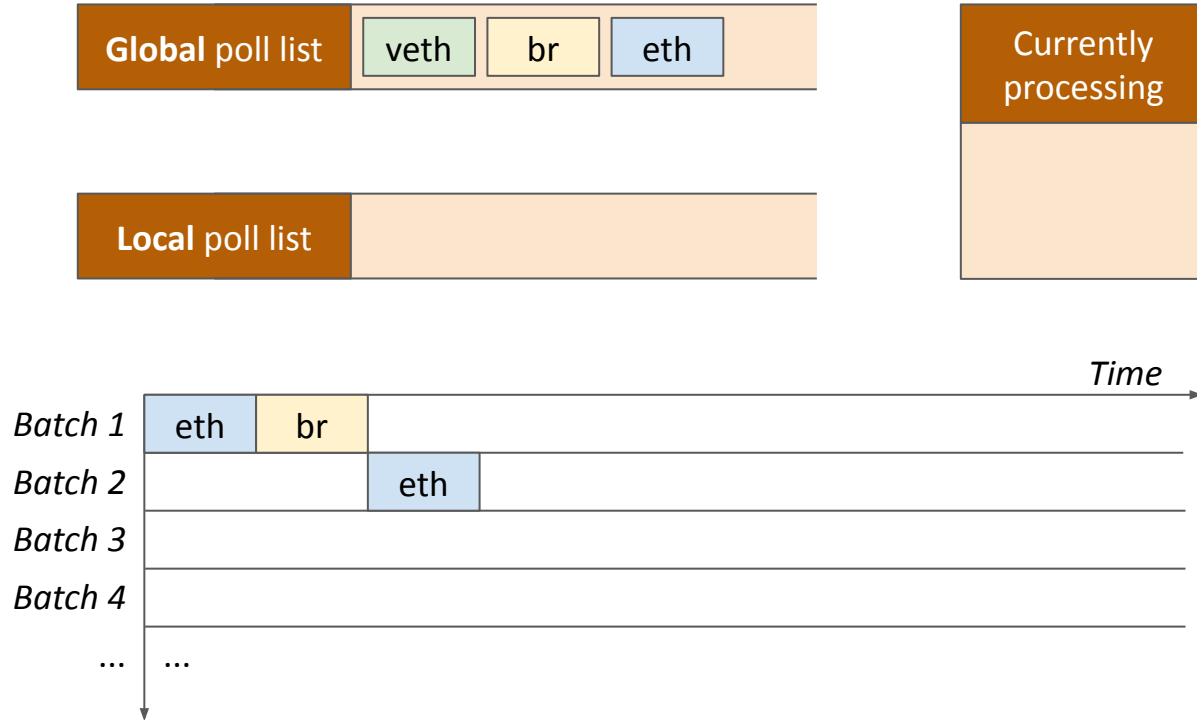
Interleaved NAPI device processing



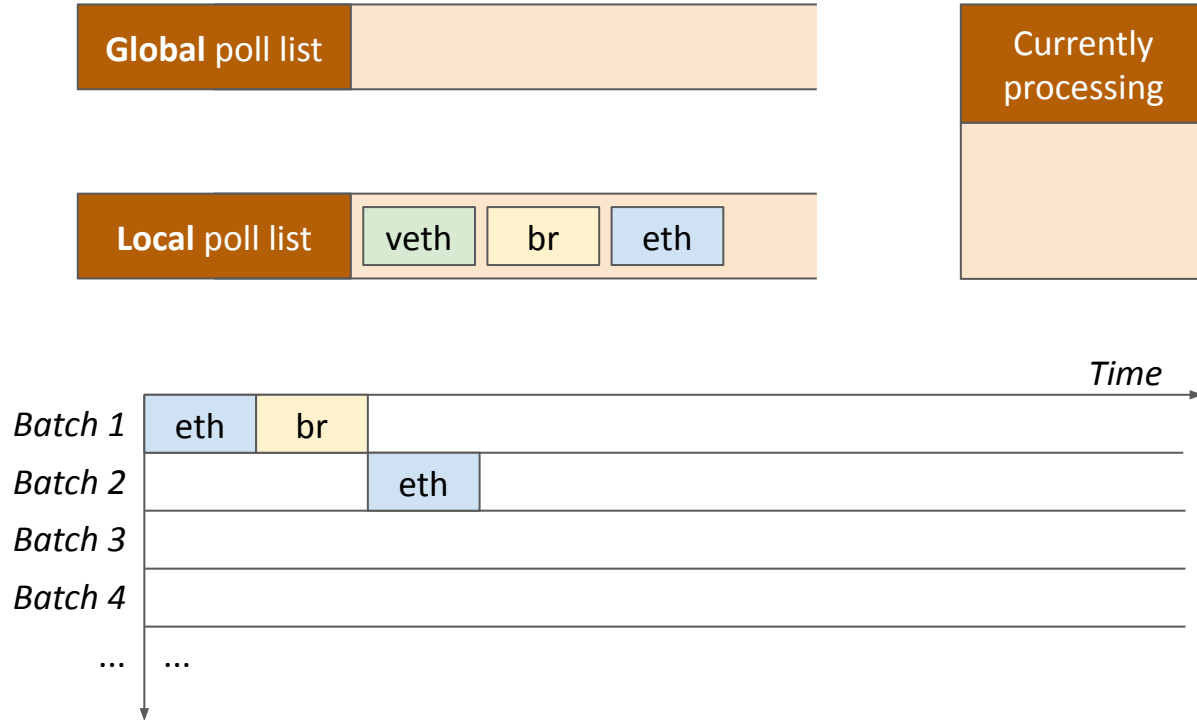
Interleaved NAPI device processing



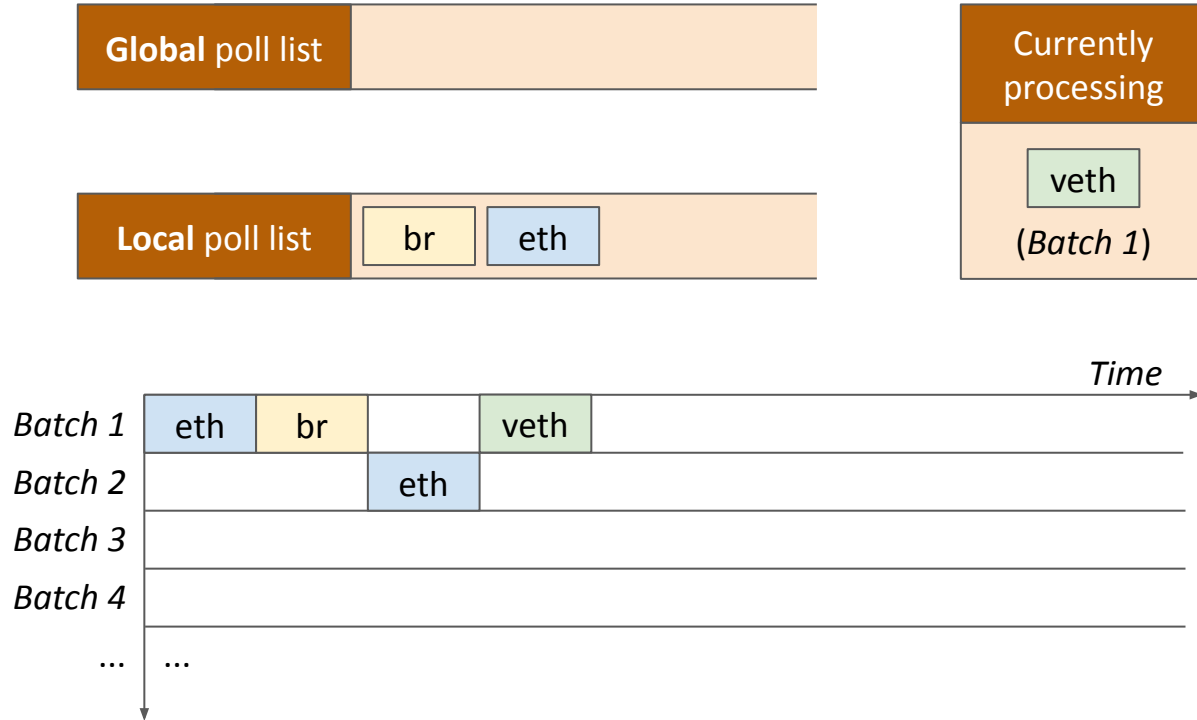
Interleaved NAPI device processing



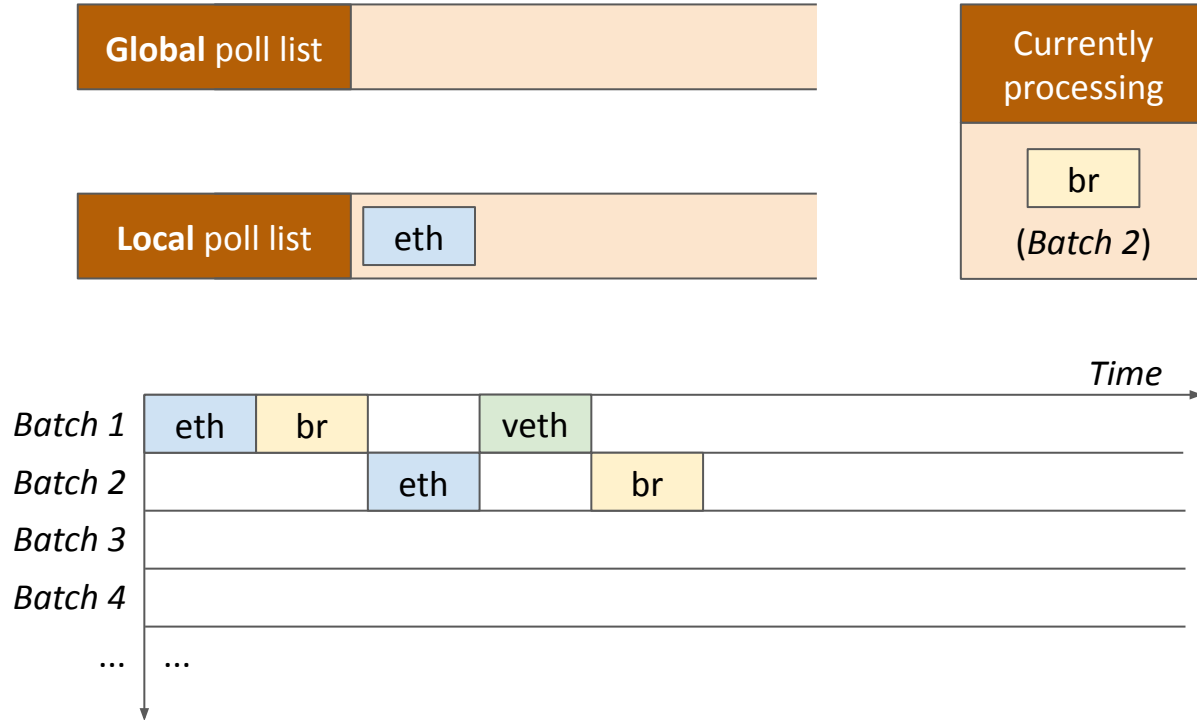
Interleaved NAPI device processing



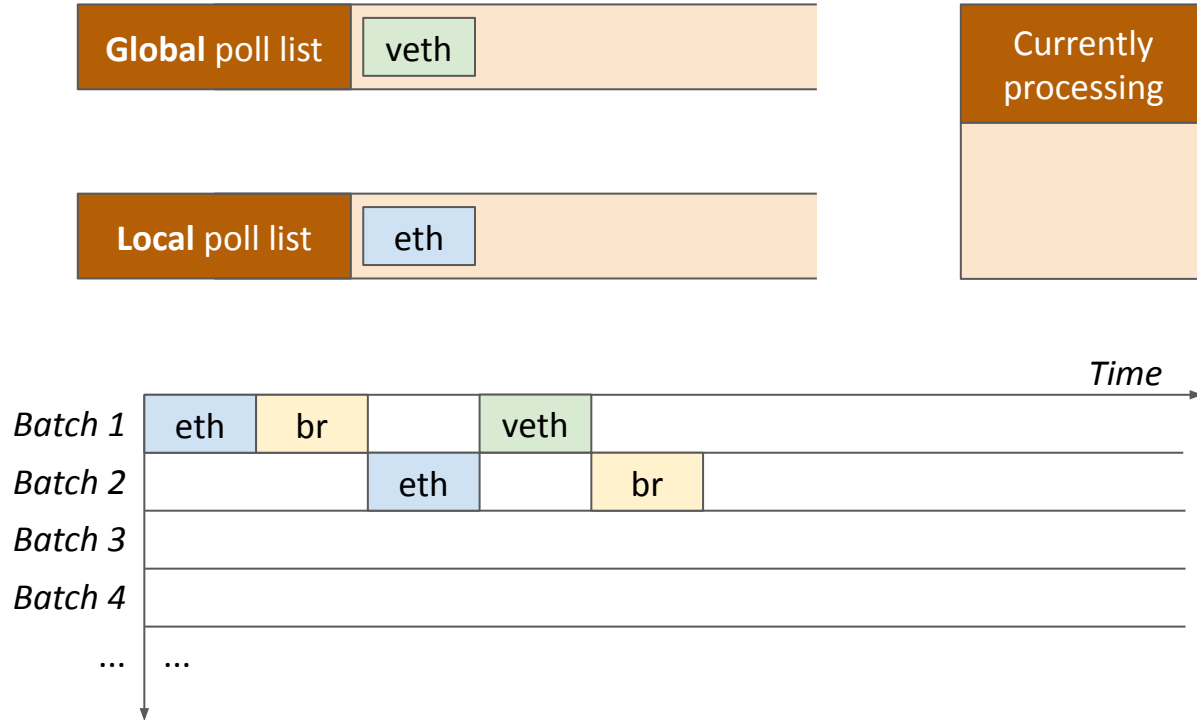
Interleaved NAPI device processing



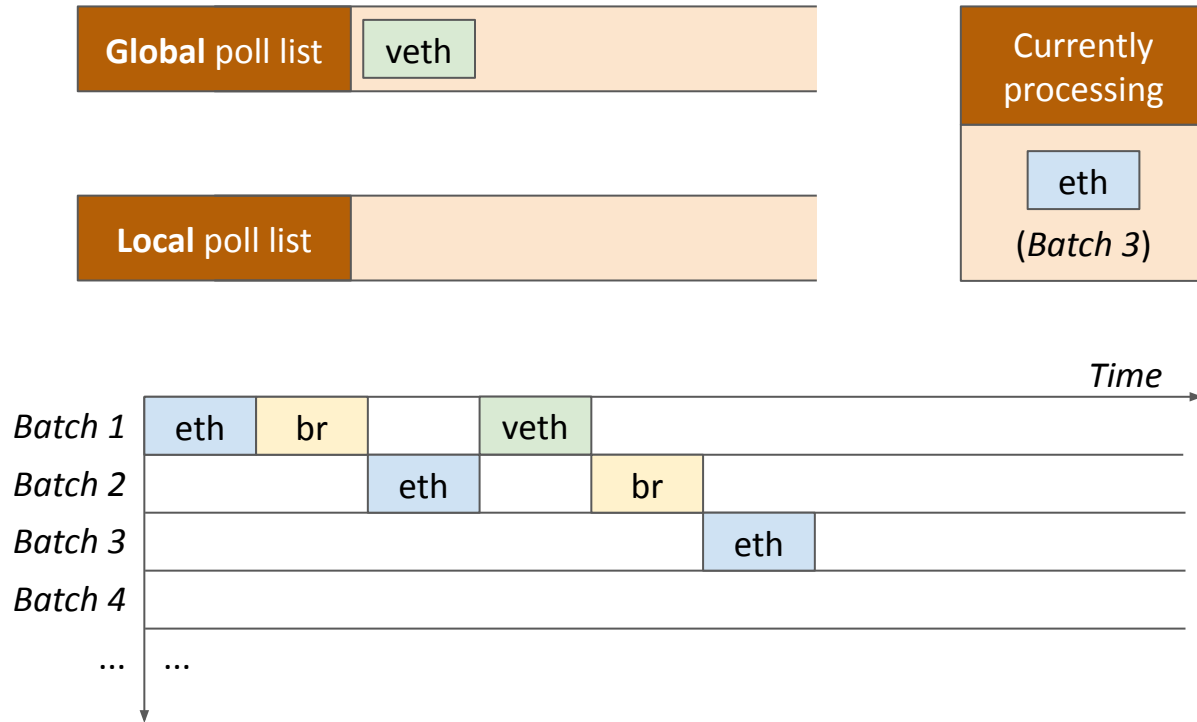
Interleaved NAPI device processing



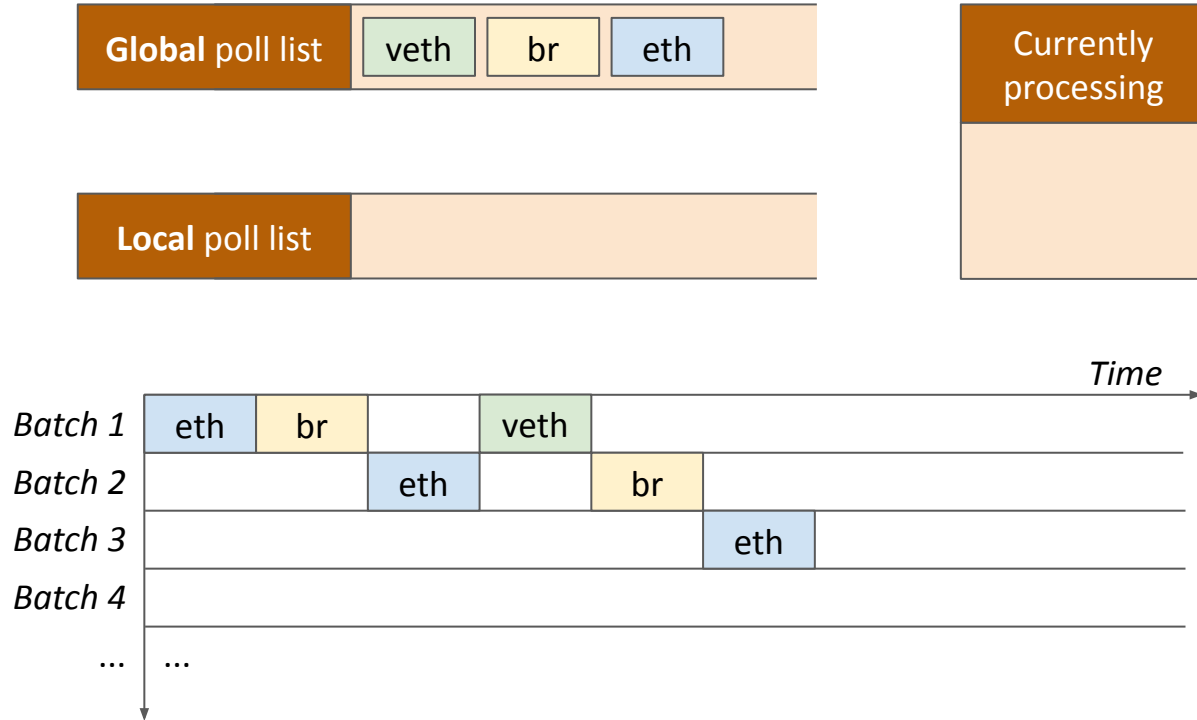
Interleaved NAPI device processing



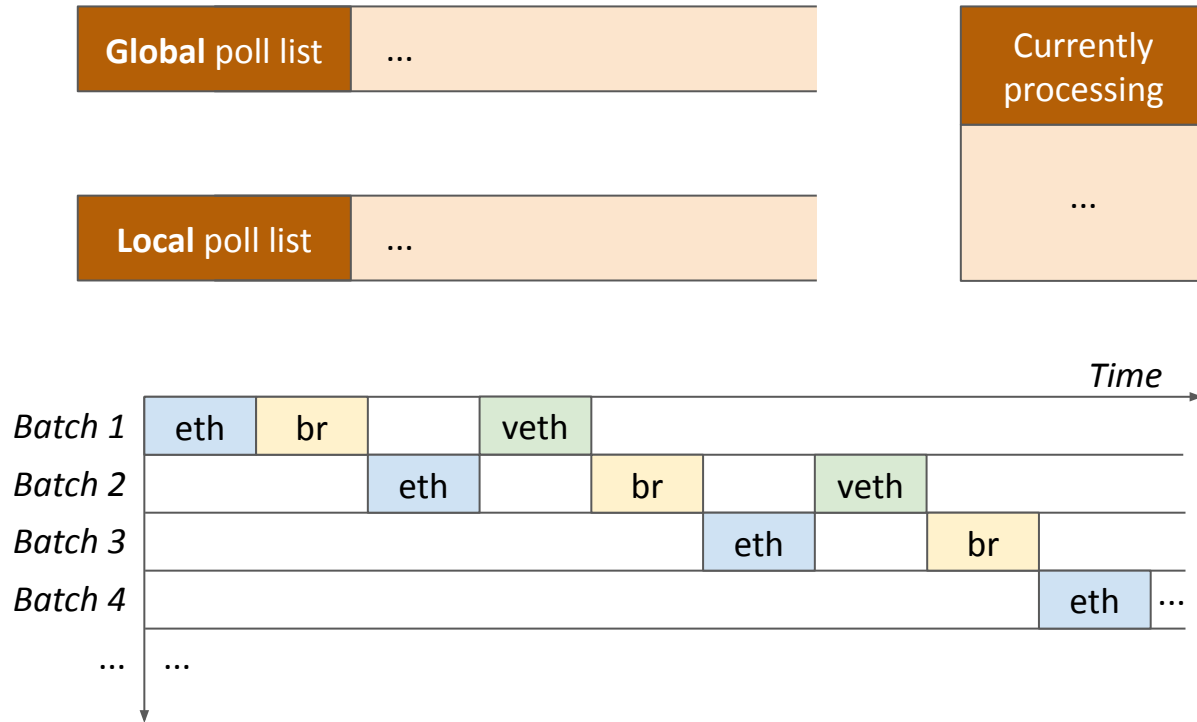
Interleaved NAPI device processing



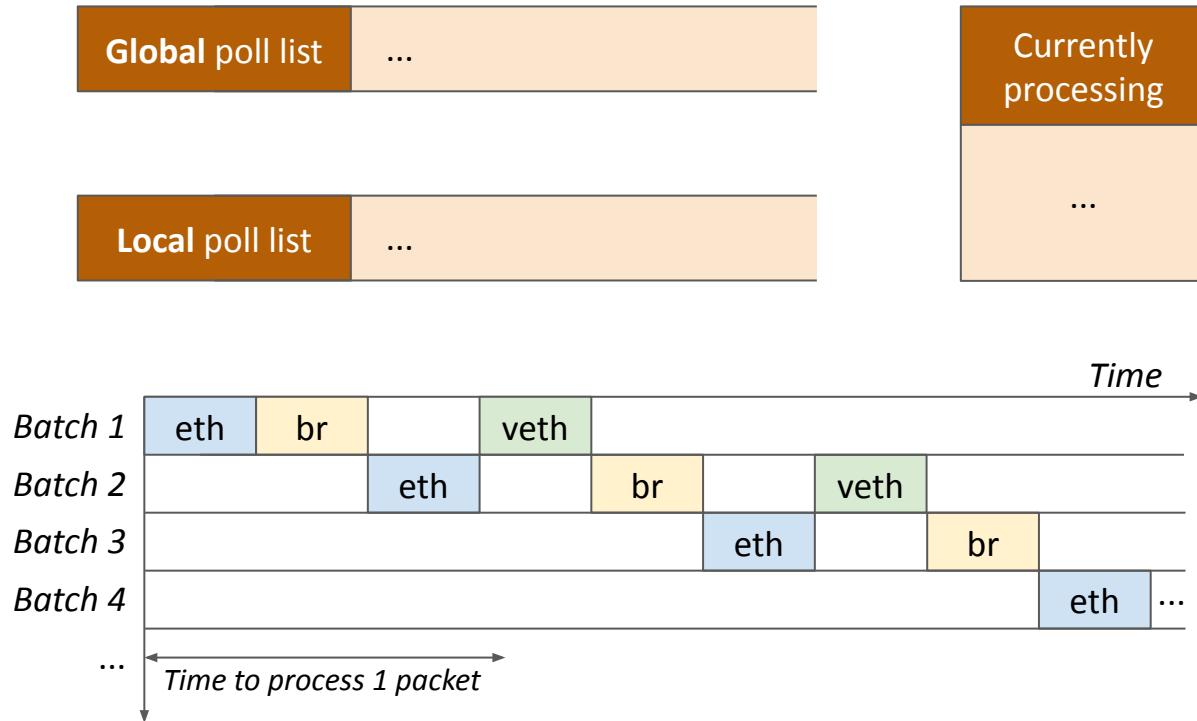
Interleaved NAPI device processing



Interleaved NAPI device processing



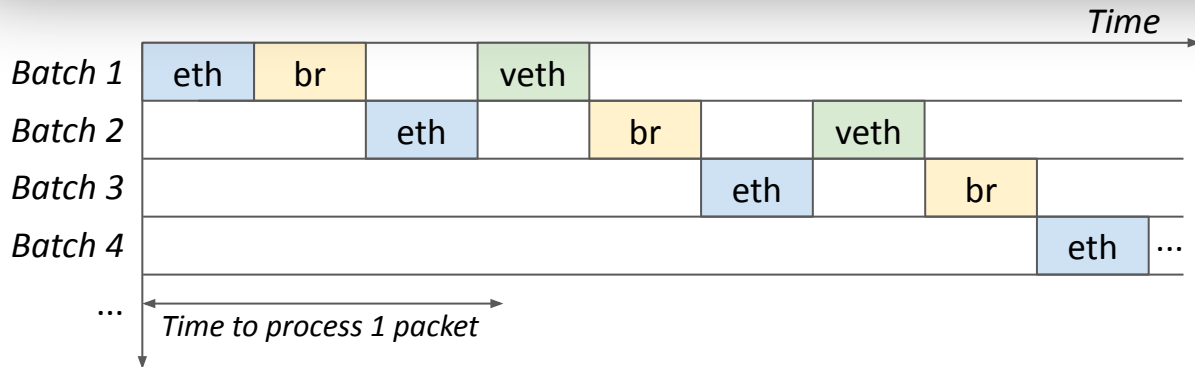
Interleaved NAPI device processing



Interleaved NAPI device processing

Summary:

Vanilla kernel network stack **hurts latency** due to **interleaved processing** stages of batches.




Lack of priority differentiation

- All packets are processed in FIFO order
- *Latency-sensitive* flows get stuck behind long queues filled with *throughput-sensitive* flows


Lack of priority differentiation

- All packets are processed in FIFO order
- *Latency-sensitive* flows get stuck behind long queues filled with *throughput-sensitive* flows




Lack of priority differentiation

- All packets are processed in FIFO order
- *Latency-sensitive* flows get stuck behind long queues filled with *throughput-sensitive* flows



Lack of priority differentiation

- All packets are processed in FIFO order
- *Latency-sensitive* flows get stuck behind long queues filled with



Is it possible to minimize latency of some flows while still offering sufficient throughput for other flows?

PRISM Design

Priority-based Streammlined Packet Processing

- Improved NAPI design
- Priority differentiation
- Streamlined NAPI device polling

PRISM Design

Priority-based Streamlined Packet Processing



- Improved NAPI design
- Priority differentiation
- Streamlined NAPI device polling

Benefits over kernel-bypass / custom network stack:



- Completely backward compatible
- No expensive hardware needed
- No need to change application code
- Preserves kernel's security, portability, and reliability

Priority differentiation



Priority differentiation




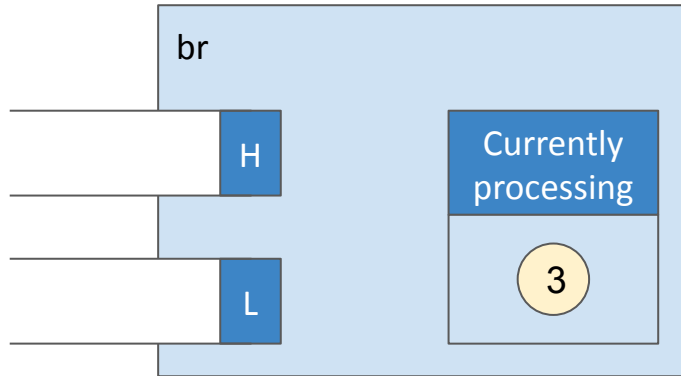
Priority differentiation




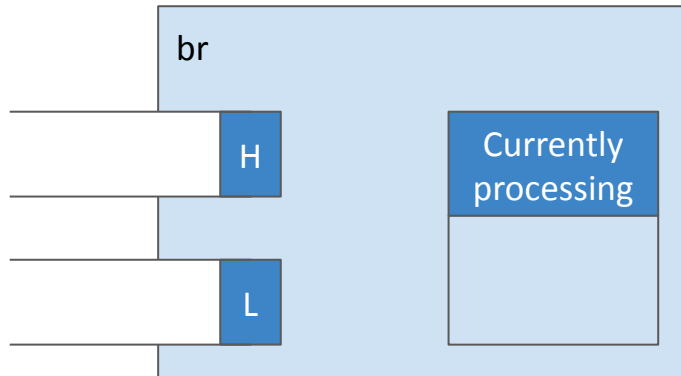
Priority differentiation




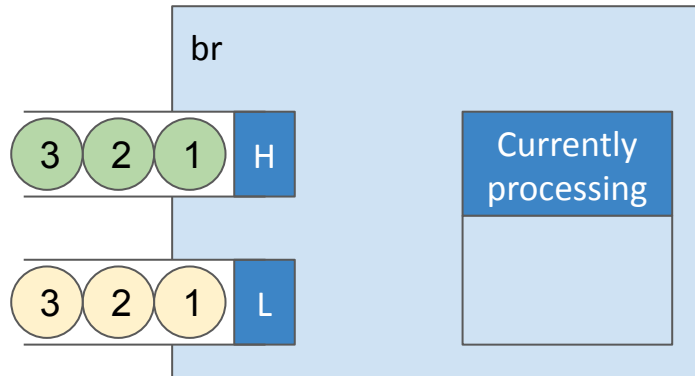
Priority differentiation




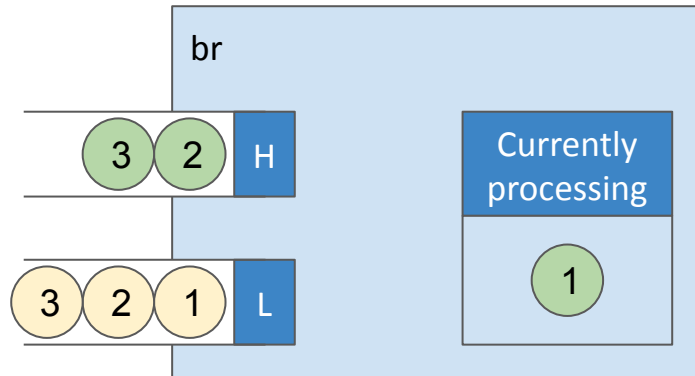
Priority differentiation




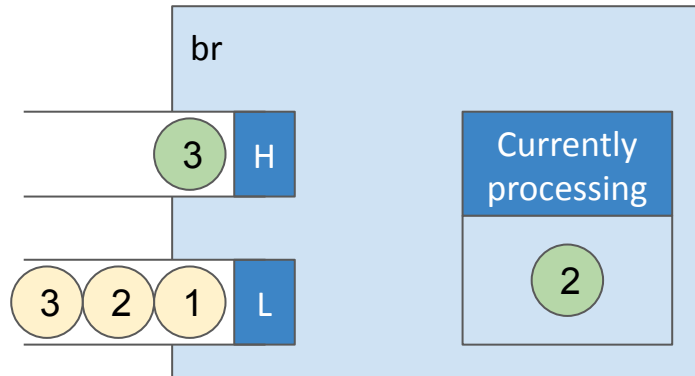
Priority differentiation




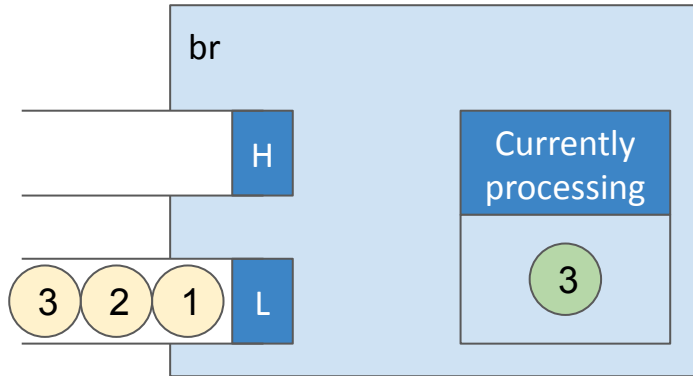
Priority differentiation




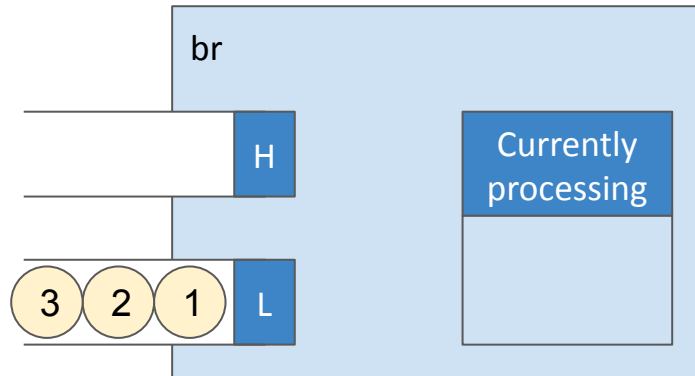
Priority differentiation



Priority differentiation



Priority differentiation



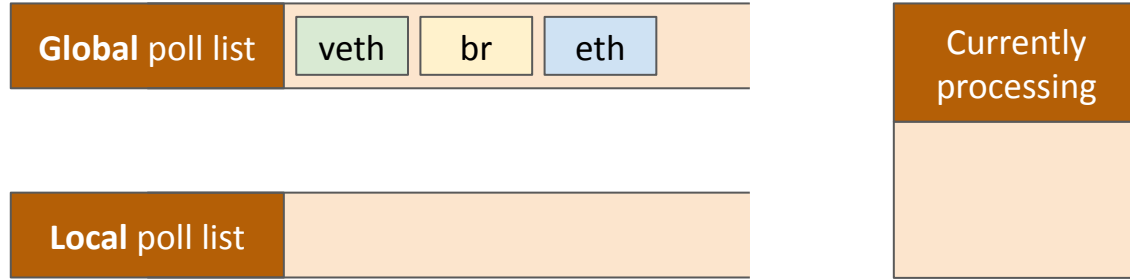
Priority differentiation

PRISM enables priority differentiation by:



- Using a dedicated high-priority packet queue
- Batch-level preemption

PRISM-batch

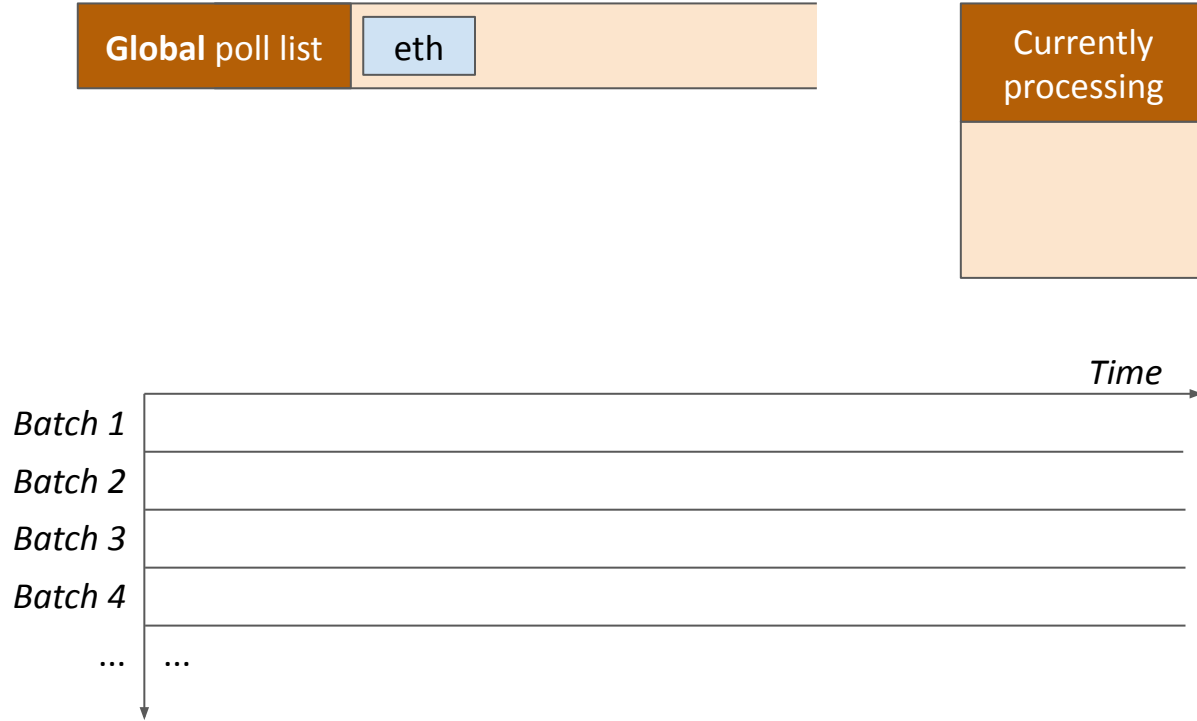
PRISM-batch



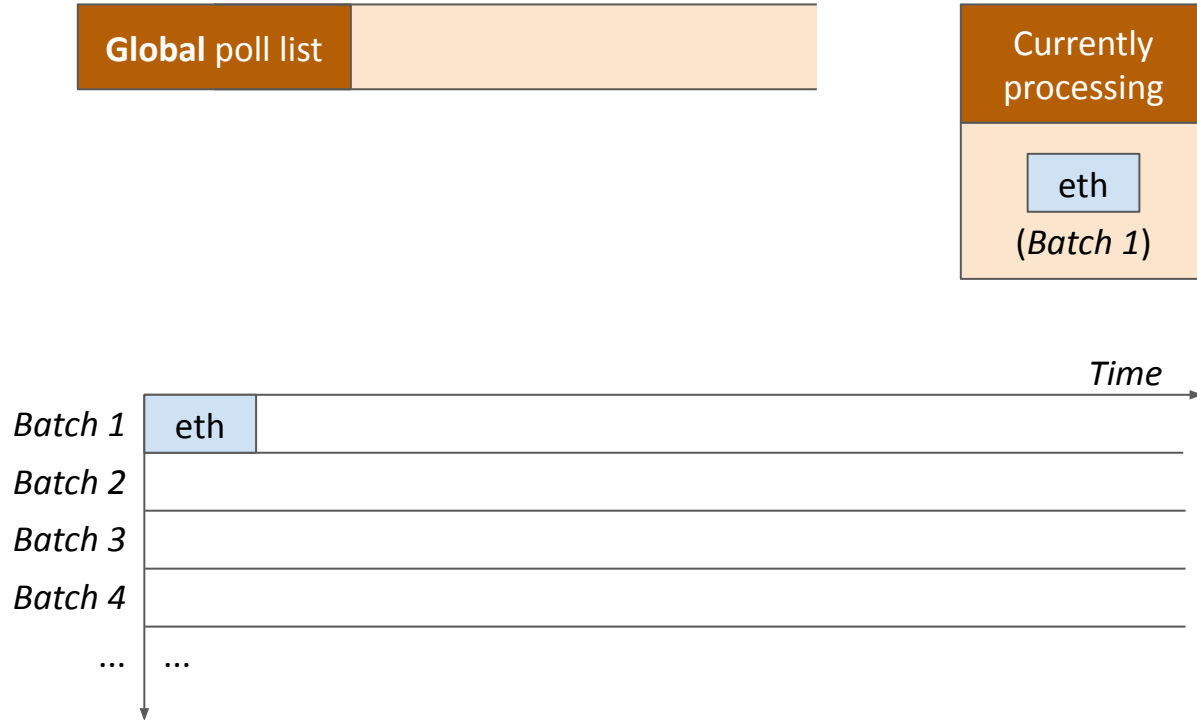
PRISM-batch



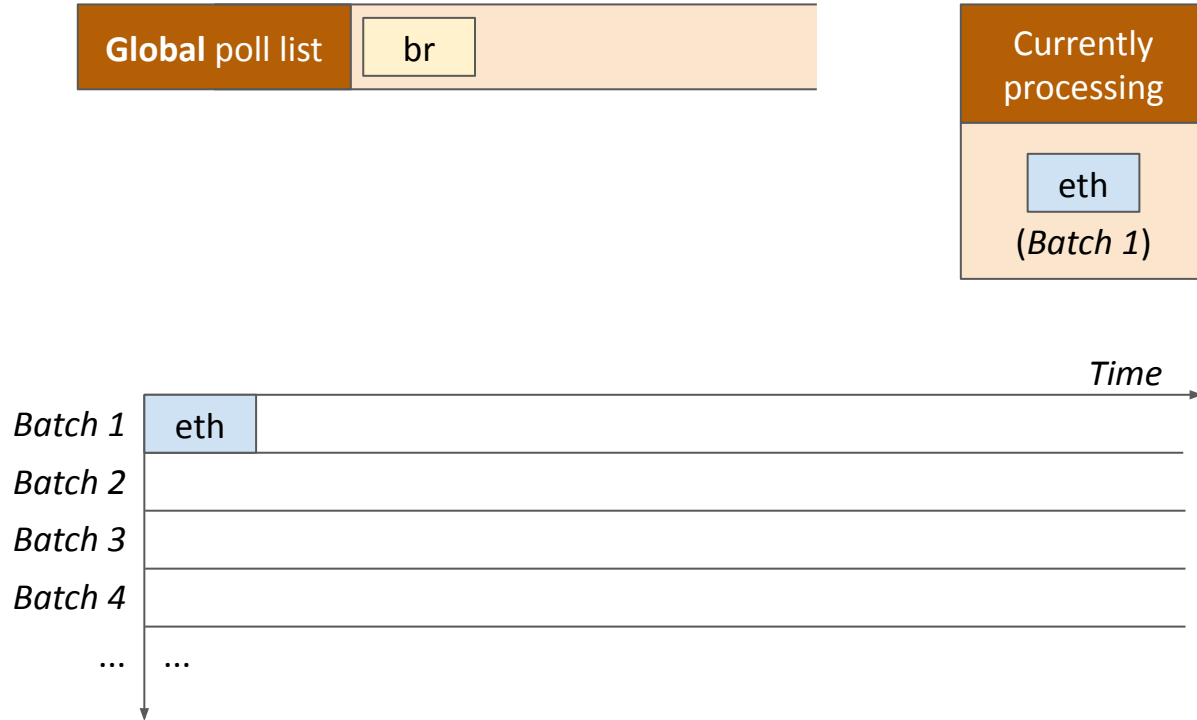
PRISM-batch



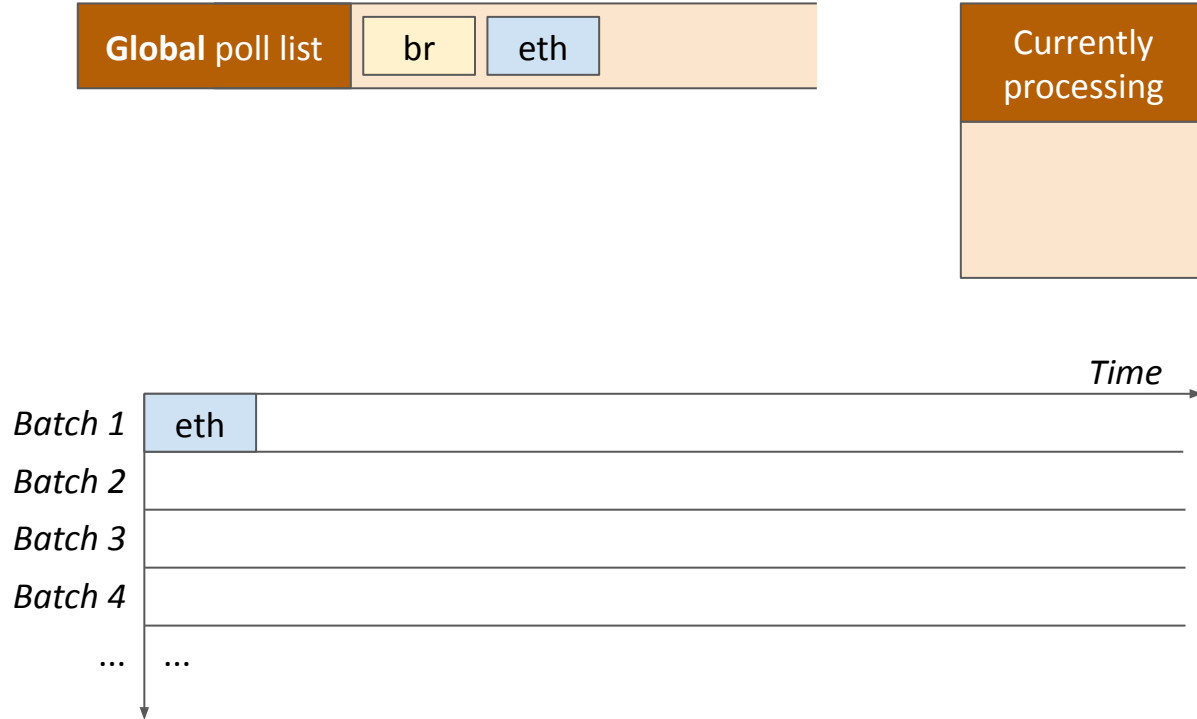
PRISM-batch



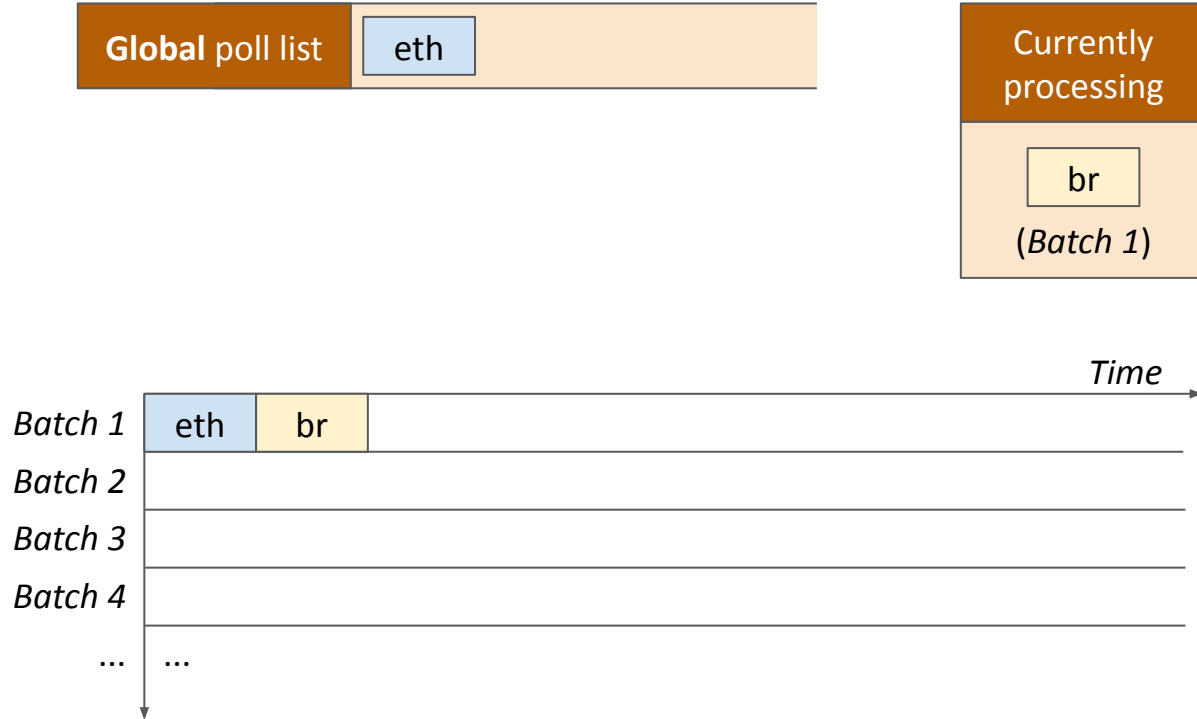
PRISM-batch



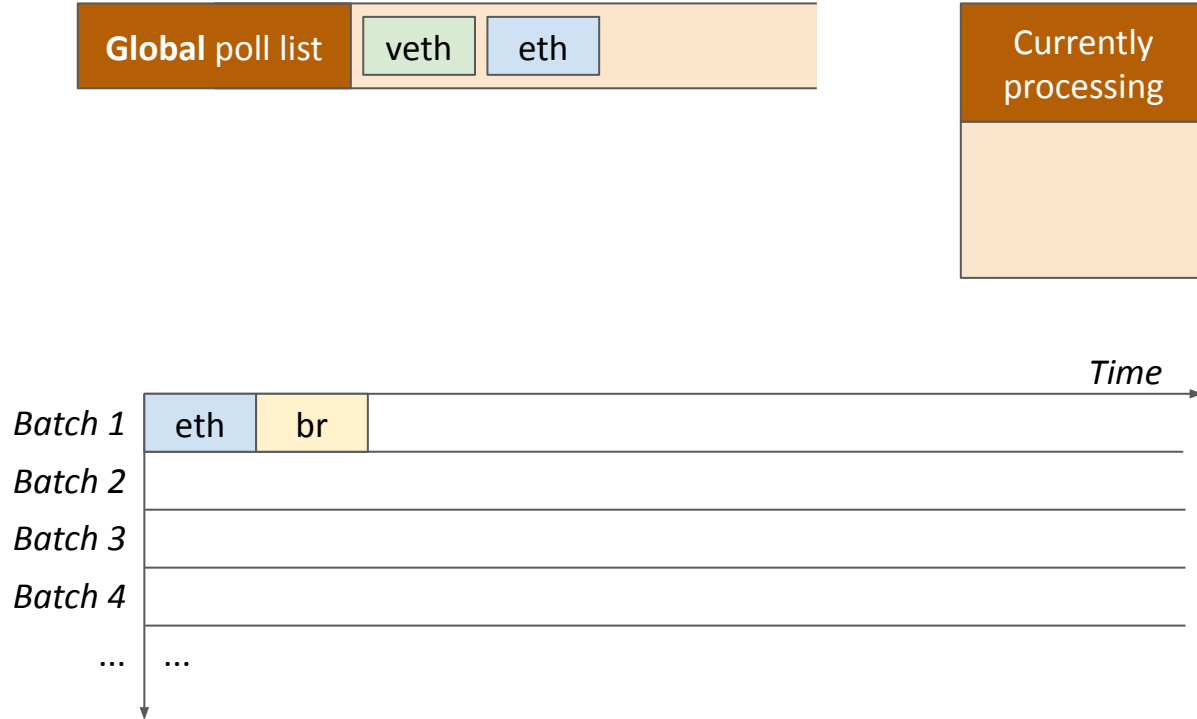
PRISM-batch



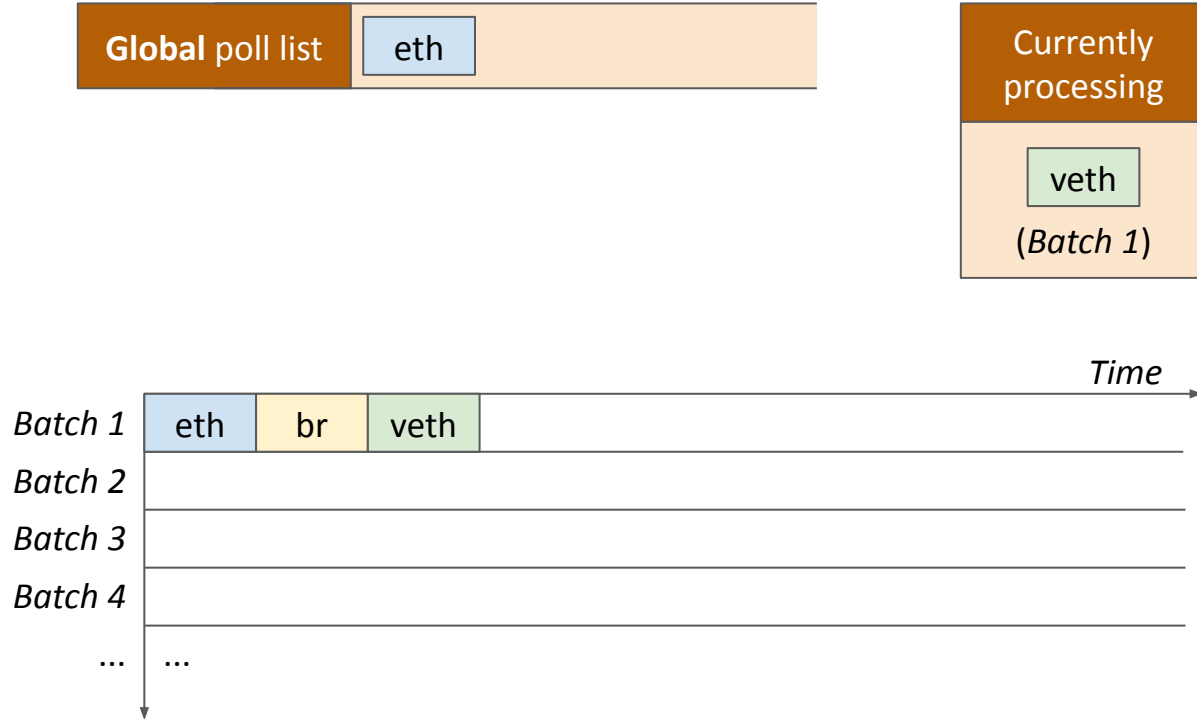
PRISM-batch



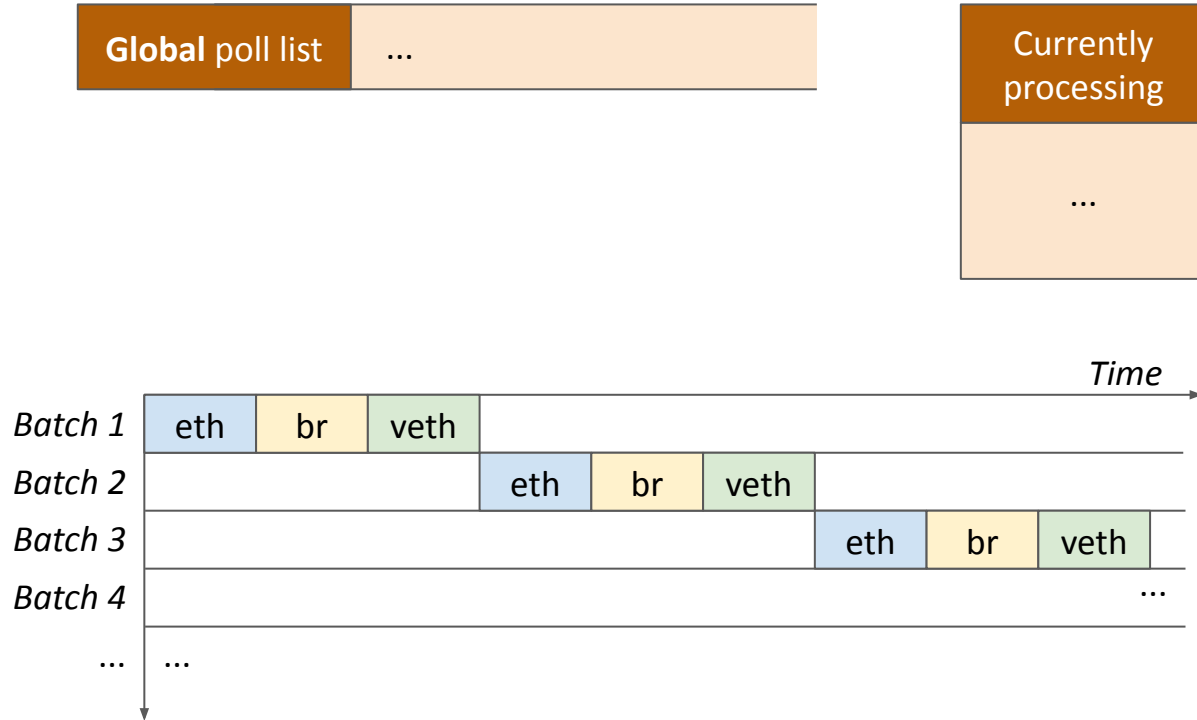
PRISM-batch



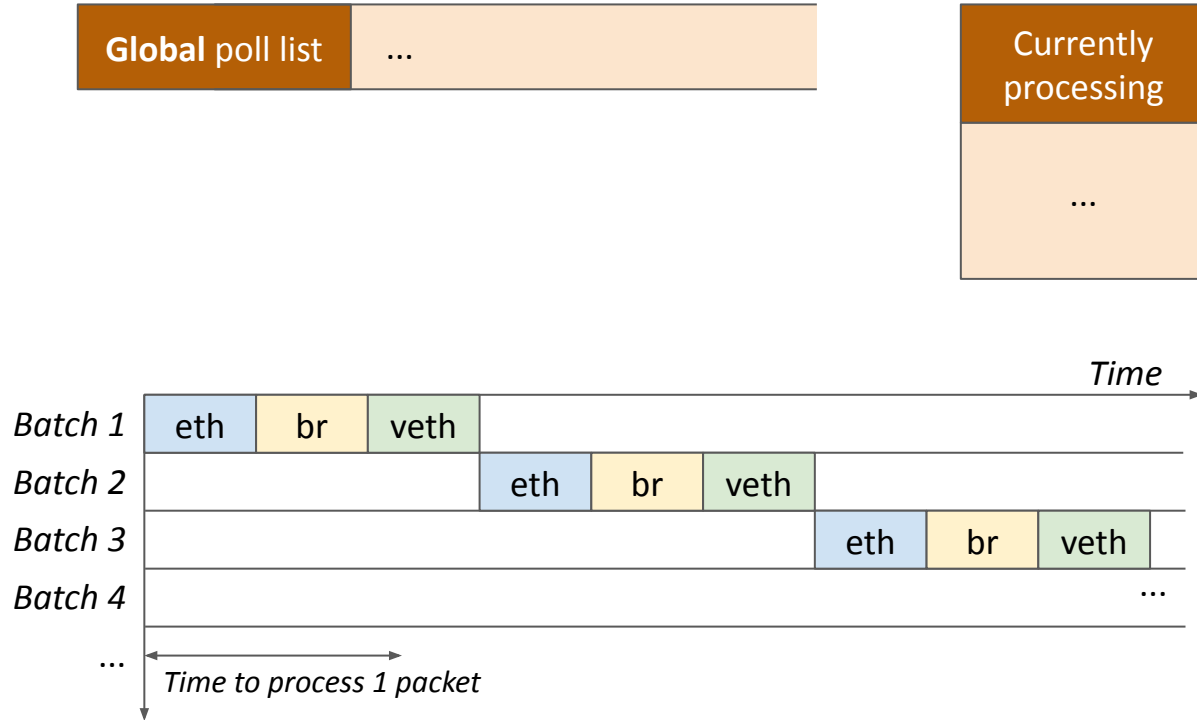
PRISM-batch



PRISM-batch



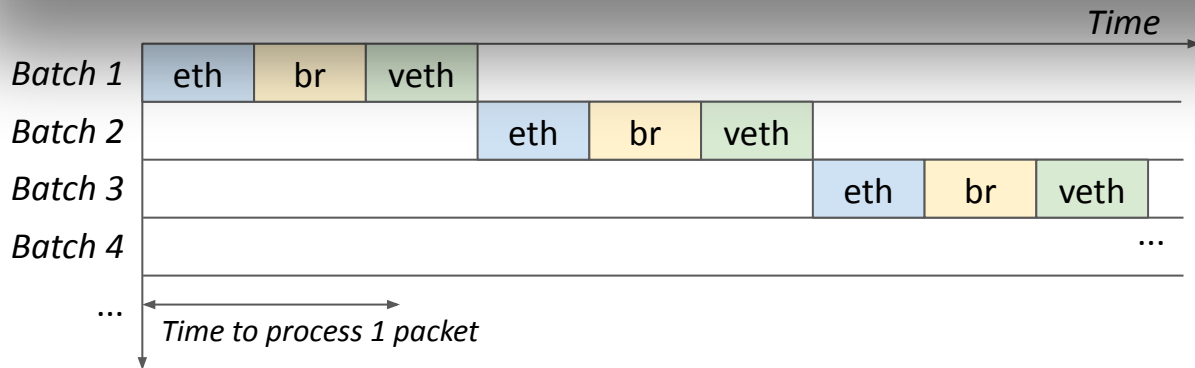
PRISM-batch



PRISM-batch

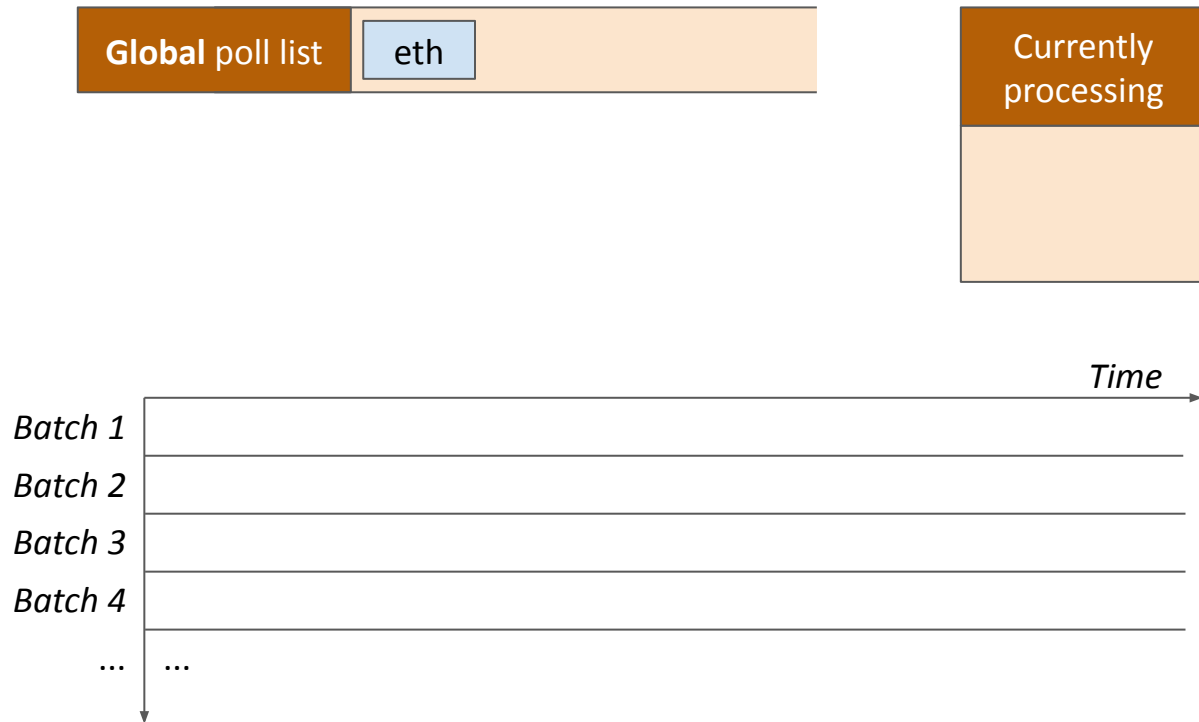
PRISM-batch:

- Improves latency by streamlining processing of batches
- Enables batch-level preemption
- But still maintains batching benefits

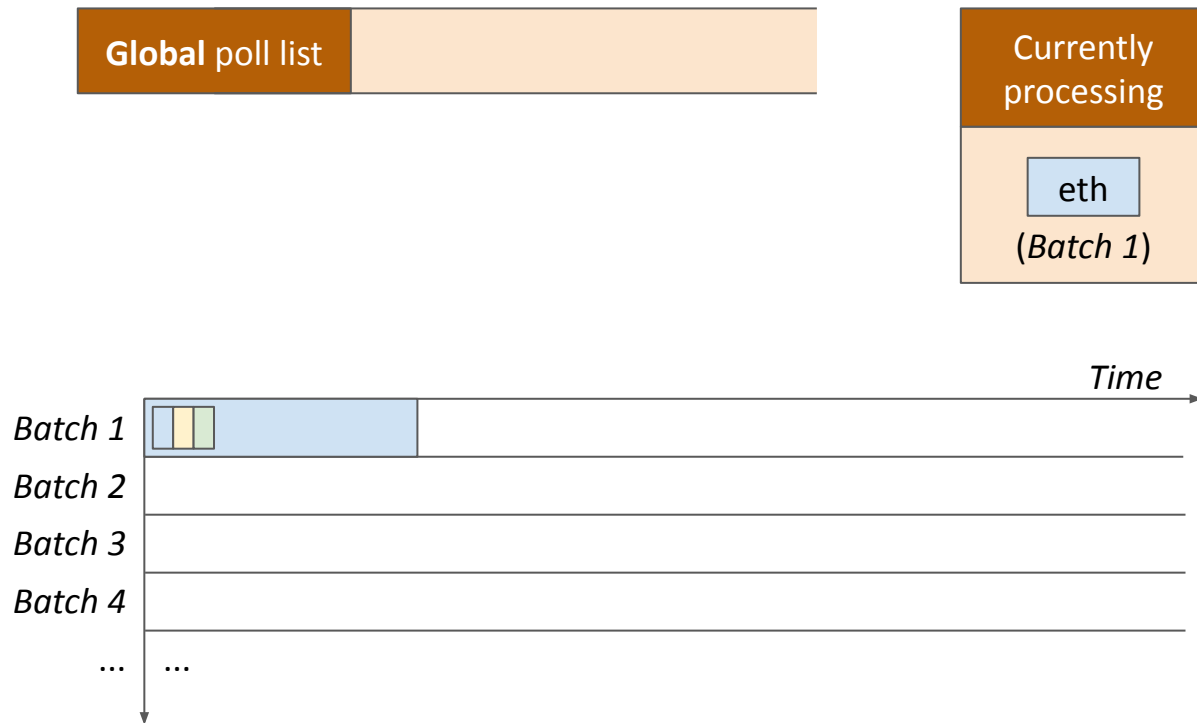


PRISM-sync

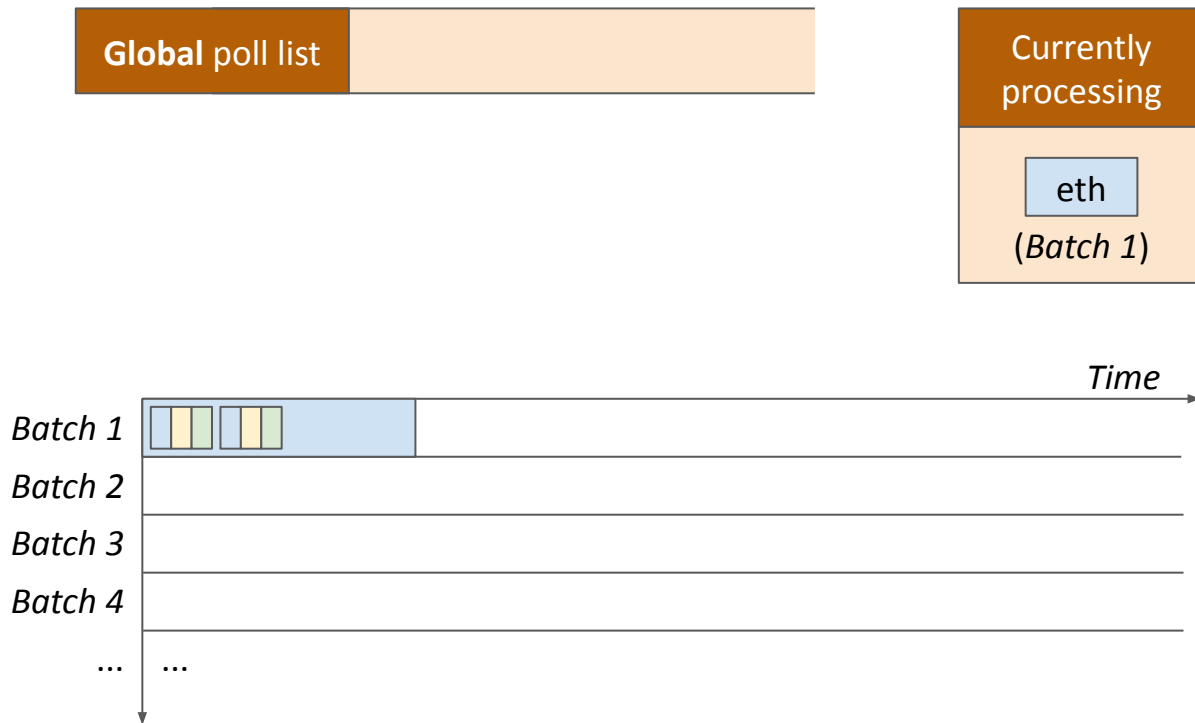
PRISM-sync



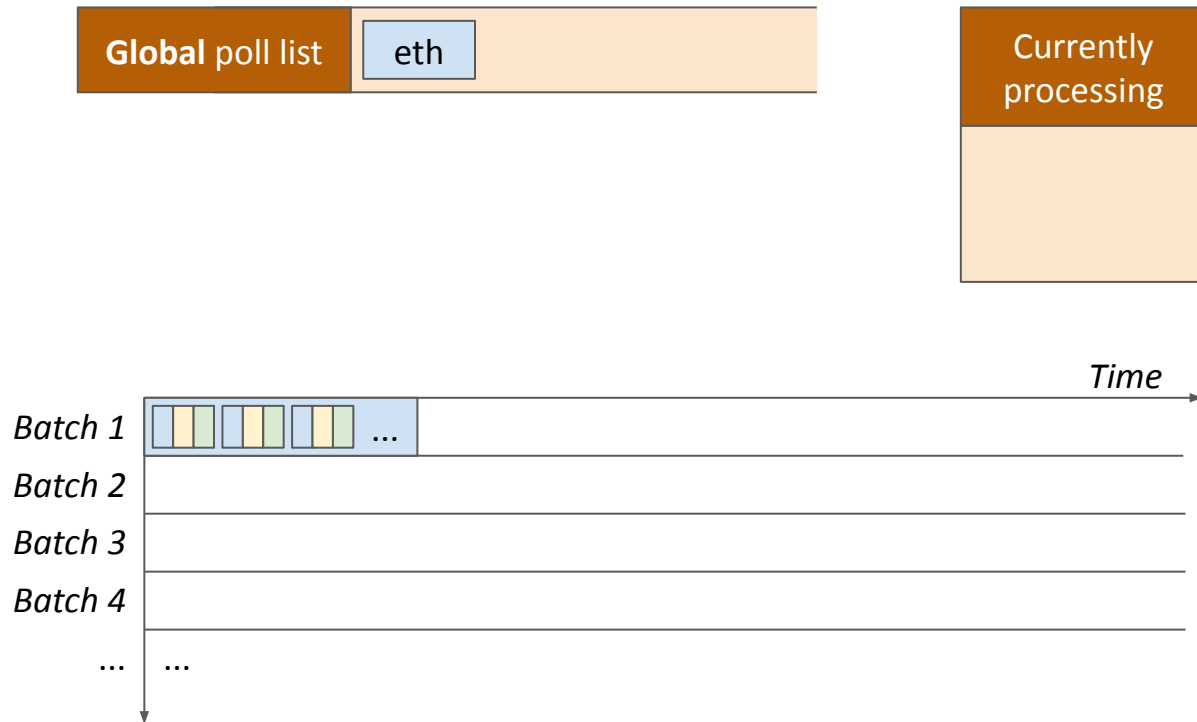
PRISM-sync



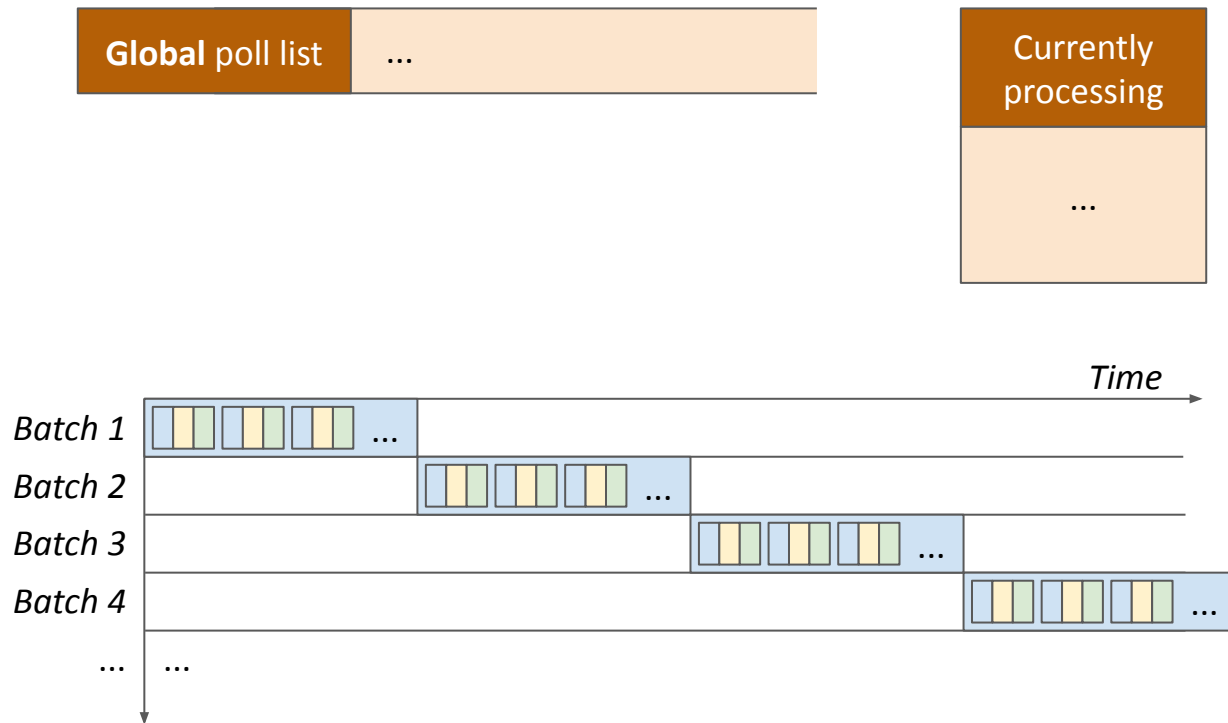
PRISM-sync



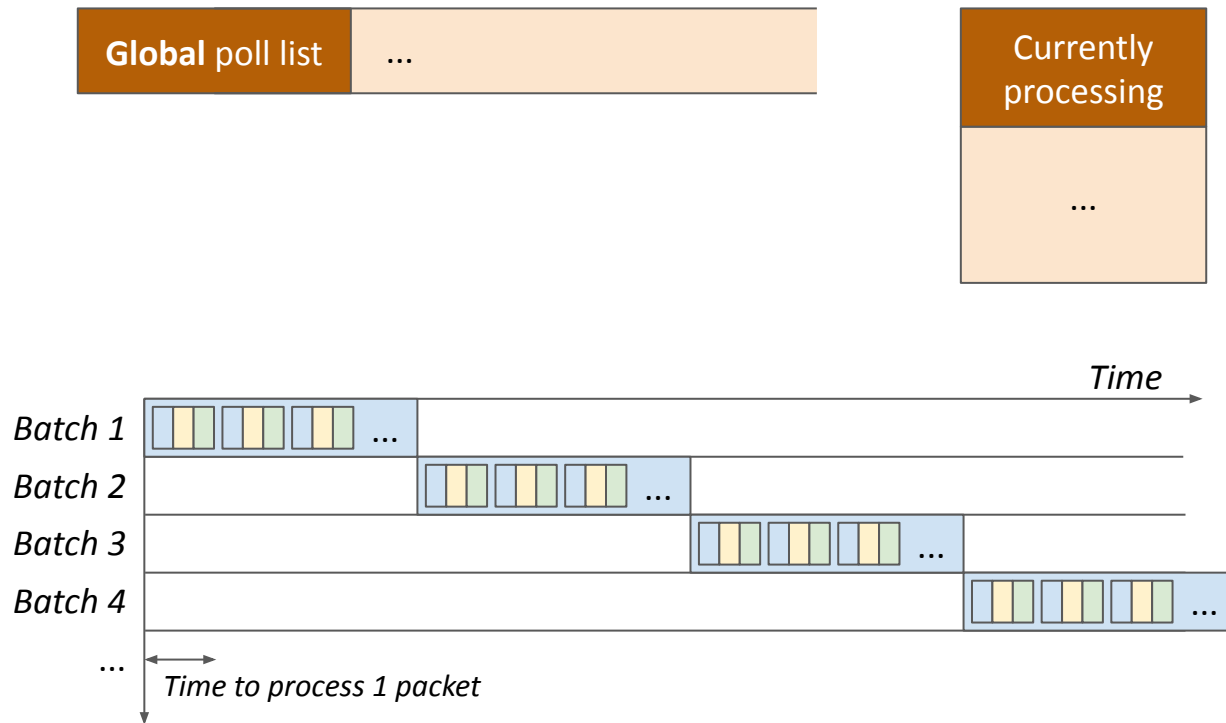
PRISM-sync



PRISM-sync




PRISM-sync



PRISM-sync

PRISM-sync:

- minimizes latency by processing each packet to completion
- loses batching benefits



Implementation

- Prototype implemented on top of Linux 5.4
- ~550 lines of code in the kernel network stack
- Publicly available at: github.com/munikarmanish/prism

Implementation

- Prototype implemented on top of Linux 5.4
- ~550 lines of code in the kernel network stack
- Publicly available at: github.com/munikarmanish/prism

Some things to note:

- Priority identification is left to user
- Prioritization in the first stage (NIC)

Evaluation — Setup

Hardware: Intel Xeon, 40 logical cores @ 2.2GHz, 128GB memory

NIC: Mellanox ConnectX-5 EN (100GbE)

Software: Ubuntu 18.04, with Linux kernel 5.4

Evaluation — Setup

Hardware: Intel Xeon, 40 logical cores @ 2.2GHz, 128GB memory

NIC: Mellanox ConnectX-5 EN (100GbE)

Software: Ubuntu 18.04, with Linux kernel 5.4

Comparison:

- Vanilla vs. PRISM-batch vs. PRISM-sync
- Idle vs. Busy server

Evaluation — Setup

Hardware: Intel Xeon, 40 logical cores @ 2.2GHz, 128GB memory

NIC: Mellanox ConnectX-5 EN (100GbE)

Software: Ubuntu 18.04, with Linux kernel 5.4



Comparison:

- Vanilla vs. PRISM-batch vs. PRISM-sync
- Idle vs. Busy server


Experiments:

- Microbenchmarks
- Application benchmarks (Memcached, Web serving)
- *[More in the paper]*


Microbenchmark — Streamlined NAPI processing





Microbenchmark — Priority differentiation




Microbenchmark — Priority differentiation



Memcached benchmark



Web serving benchmark



Conclusion

- Linux kernel has many sources of inefficiencies
 - Especially for overlay packet processing
 - Interleaved packet processing logic
 - No support for flow prioritization
- PRISM improves the performance of container flows
 - **Priority-based** packet queue management
 - **Streamlined** multi-stage processing pipeline
 - **Batch-level preemption** or synchronous processing
- Implementation: github.com/munikarmanish/prism

Conclusion

- Linux kernel has many sources of inefficiencies
 - Especially for overlay packet processing
 - Interleaved packet processing logic
 - No support for flow prioritization
- PRISM improves the performance of container flows
 - **Priority-based** packet queue management
 - **Streamlined** multi-stage processing pipeline
 - **Batch-level preemption** or synchronous processing
- Implementation: github.com/munikarmanish/prism

For more info, please contact me at manish.munika@uta.edu.

Backup slides

Background: NAPI

- Interrupt-based packet processing
 - IRQ raised for each packet received by the NIC
 - Good for light load, bad for heavy load
- NAPI
 - Many packets are processed (in batches) after a single interrupt
 - Interrupt-mode for light load
 - Poll-mode for heavy load
- RSS / RPS / RFS
 - Distributes independent flows to separate cores for processing
- Still naive FIFO queueing!
 - High-priority packets can get stuck behind long queues of low-priority packets

Remaining issues

- Flow prioritization in the NIC queue
- Priority synchronization between user application and network stack
- Multiple priority levels

Why this problem?

- NICs are becoming faster quicker
- CPUs cannot keep up
- OS packet processing logic is becoming performance bottleneck
- Moreover, containers are everywhere
- Containers use container overlay networks
- Container network incurs high overheads
- In a highly utilized system, the high priority messages experience large queueing delays

Microbenchmark — Priority differentiation (host network)

